



# RICE

Unconventional Wisdom



# OCCA: Portability Layer for Many-core Thread Programming

Rice Oil & Gas Workshop 2015  
March 3, 2015.

Tim Warburton

Department of Computational & Applied Mathematics @ Rice University



# Extended Research Team

I am fortunate to work with a team of excellent researchers



JF Remacle  
Meshing & Numerics  
Sabbatical @ Rice



Amik St-Cyr  
Royal Dutch Shell  
Adjunct Assoc Prof



Jesse Chan  
Advanced Numerics  
Post-doc



Axel Modave  
Accelerate Numerics  
Industrial Post-doc



Florian Kummer  
Multi-phase Flows  
Visiting scholar



Ali Karakus  
Two Phase Flows  
Visiting student



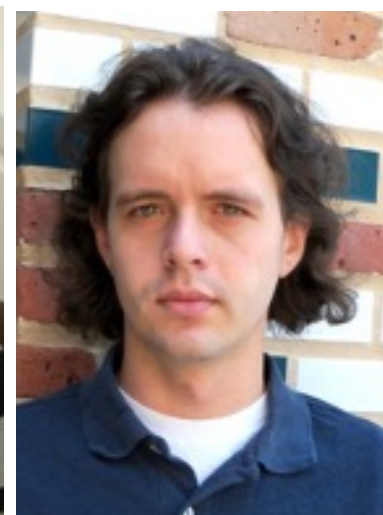
Nichole Stilwell  
CFD  
CAAM Grad > USAF



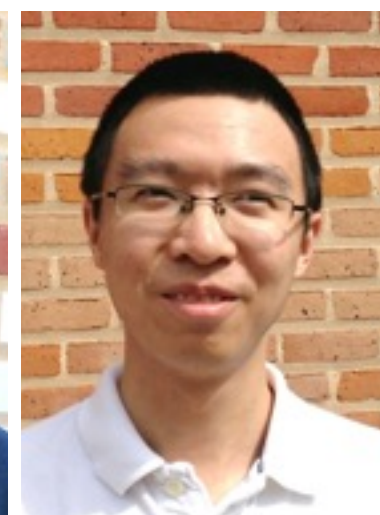
Rajesh Gandham  
Oceans & AMG  
CAAM Grad



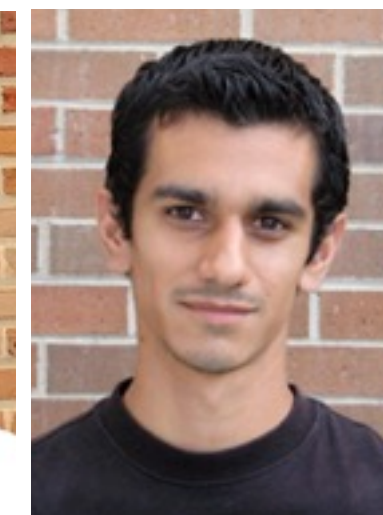
David Medina  
HPC & Accelerators  
CAAM Grad



Reid Atcheson  
Adaptivity/Helmholtz  
CAAM Grad



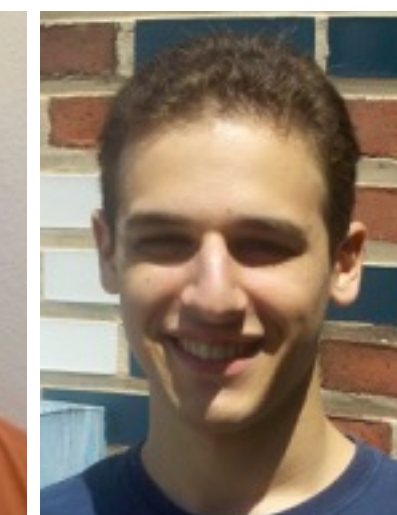
Zheng (Frank) Wang  
Radiation Transport  
CAAM Grad



Arturo Vargas  
Hermite+DG  
CAAM Grad



Tim Moon  
Bioheat modeling  
Ugrad => Stanford



Michael Frano  
Bioheat modeling  
Ugrad => BP

Industrial internships, projects & fellowships:  
Shell, BP, Halliburton, Hess, Stoneridge Technology, Hypercomp, Z-terra, ExxonMobil\*

# Overview

## Interlude on Many-core Computing:

- Challenges of modern computing: uncertainty, fragmentation, parallel algorithms.
- OCCA: unified programming model for multi- and many-core computing.

## OCCA:

- Introduce API and kernel languages, OKL and OFL
- Mention applications and benchmarks with performance comparisons

## Live Demo

- Download and install OCCA from scratch
- Display features in the programming interface and kernel languages

## Interactive Demo

- We let you (the attendees) get hands-on experience using OCCA
- Members in our research group will walk around to help if needed

# OCCA: portable many-core library

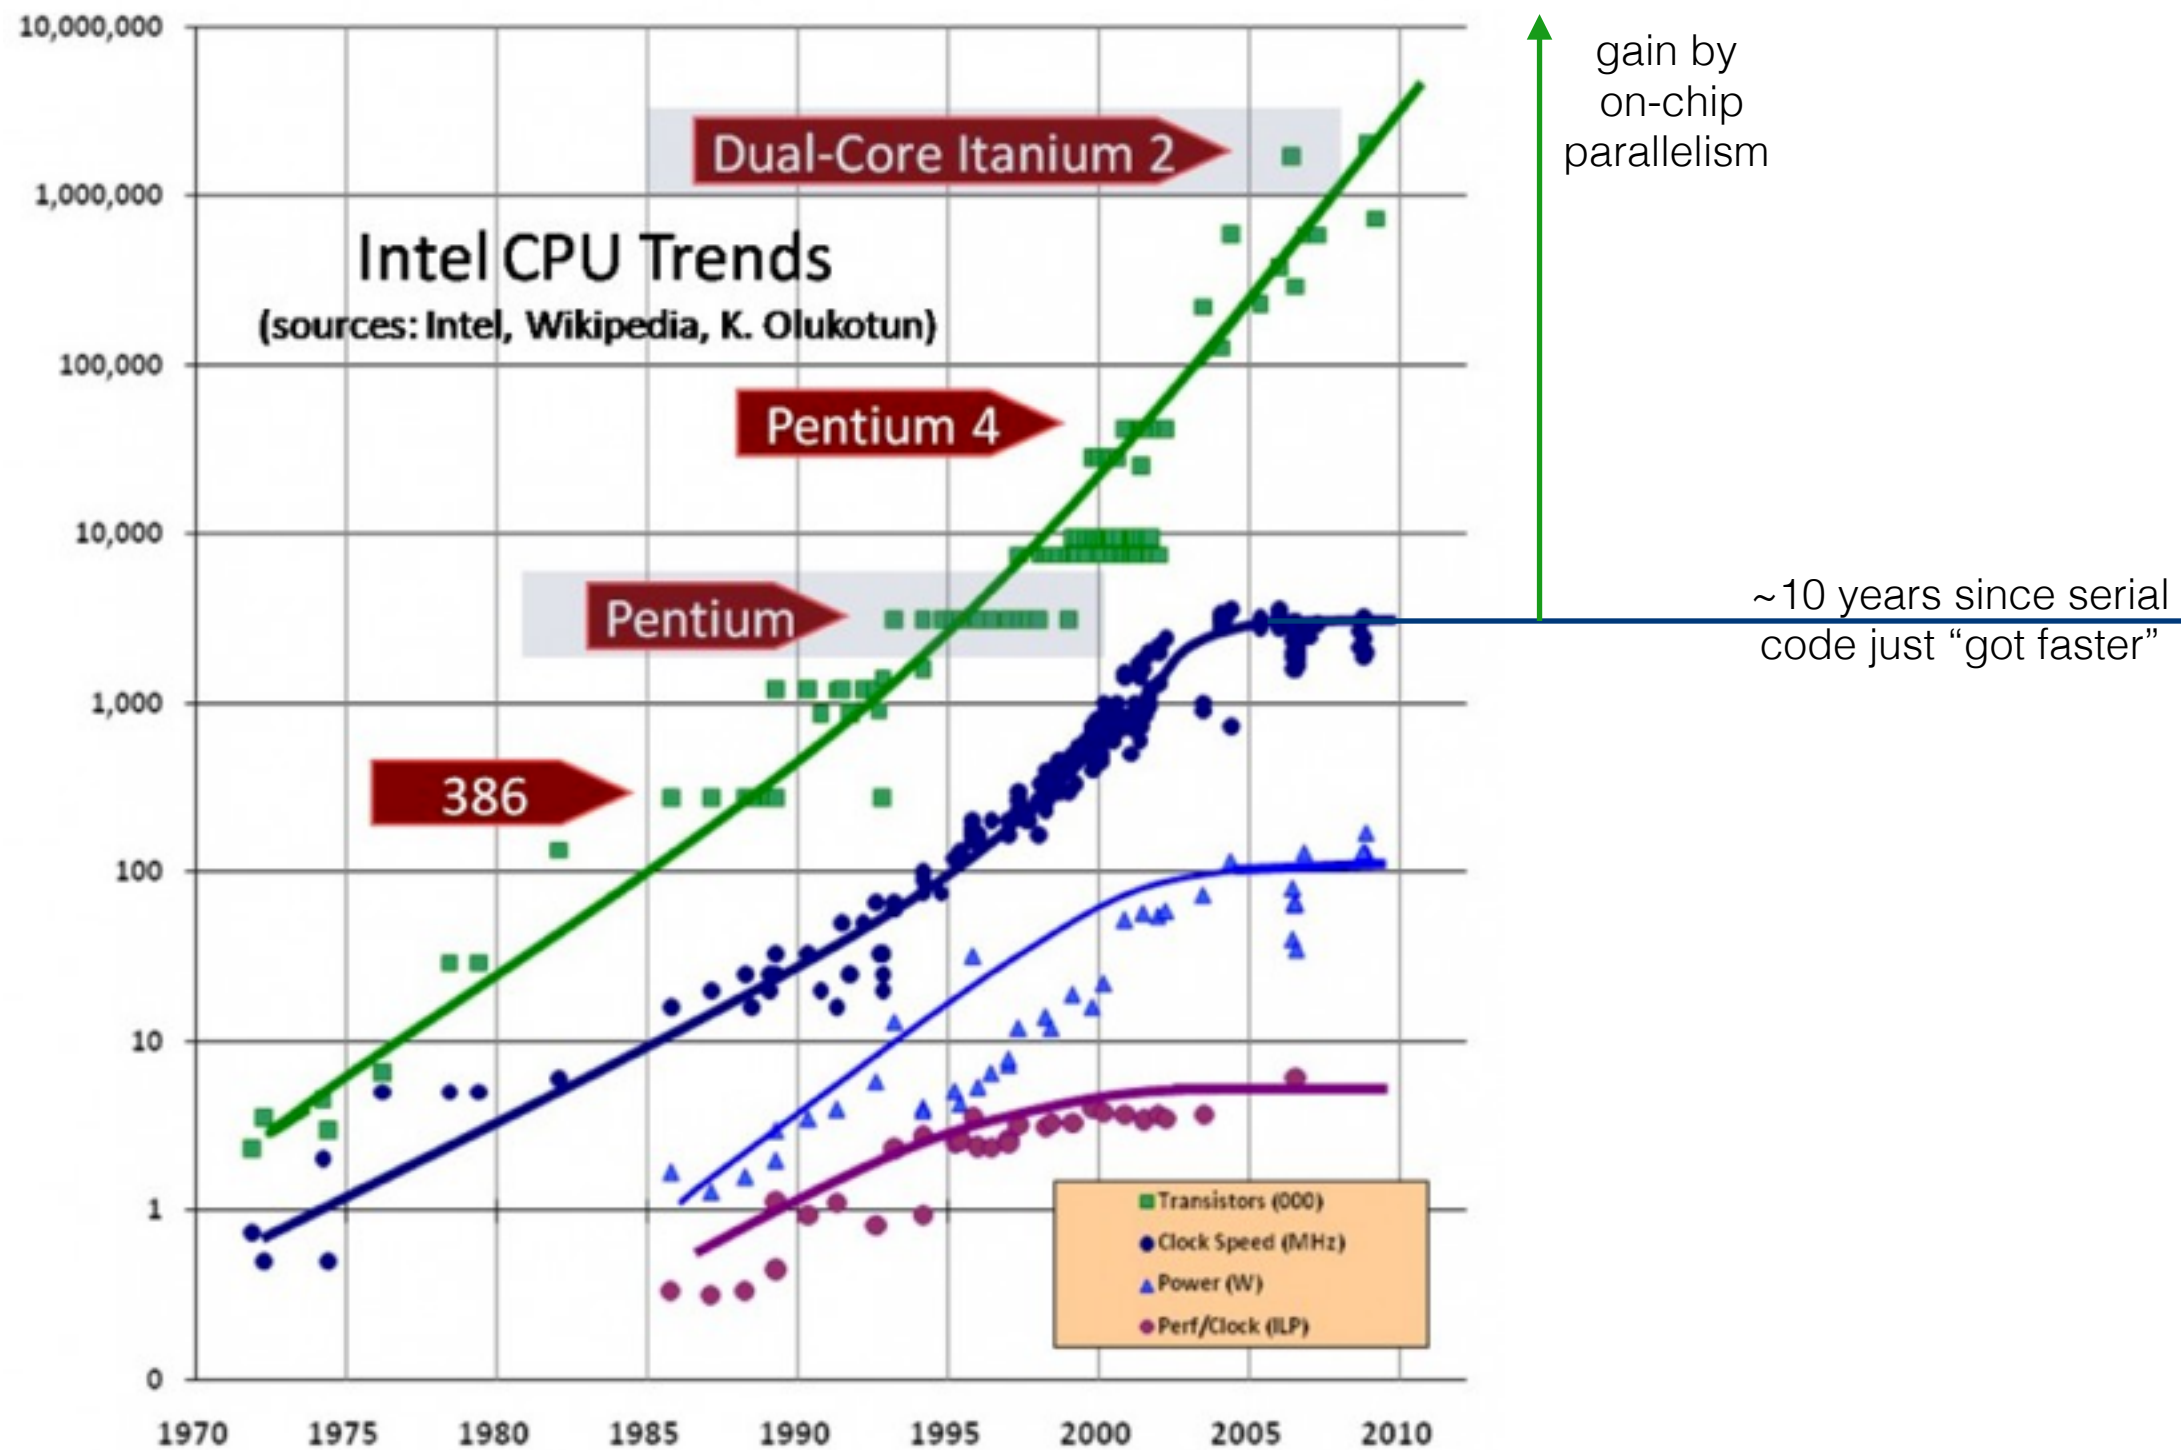
Project with David Medina (Rice) & contributors.

Funded by the Argonne CESAR Exascale Co-design center



# Moore's Law: transition to many-core

There is no escaping parallel computing any more even on a laptop.



# Ubiquitous Multicore: no escape

Multicore and heterogeneous compute is prevalent from Intel's Broadwell-U mobile processor to the upcoming Intel Knights Landing accelerator.

Broadwell die shot  
(removed)

Broadwell-U: 14nm process with dual cores  
& integrated GPU [1.9B transistors]

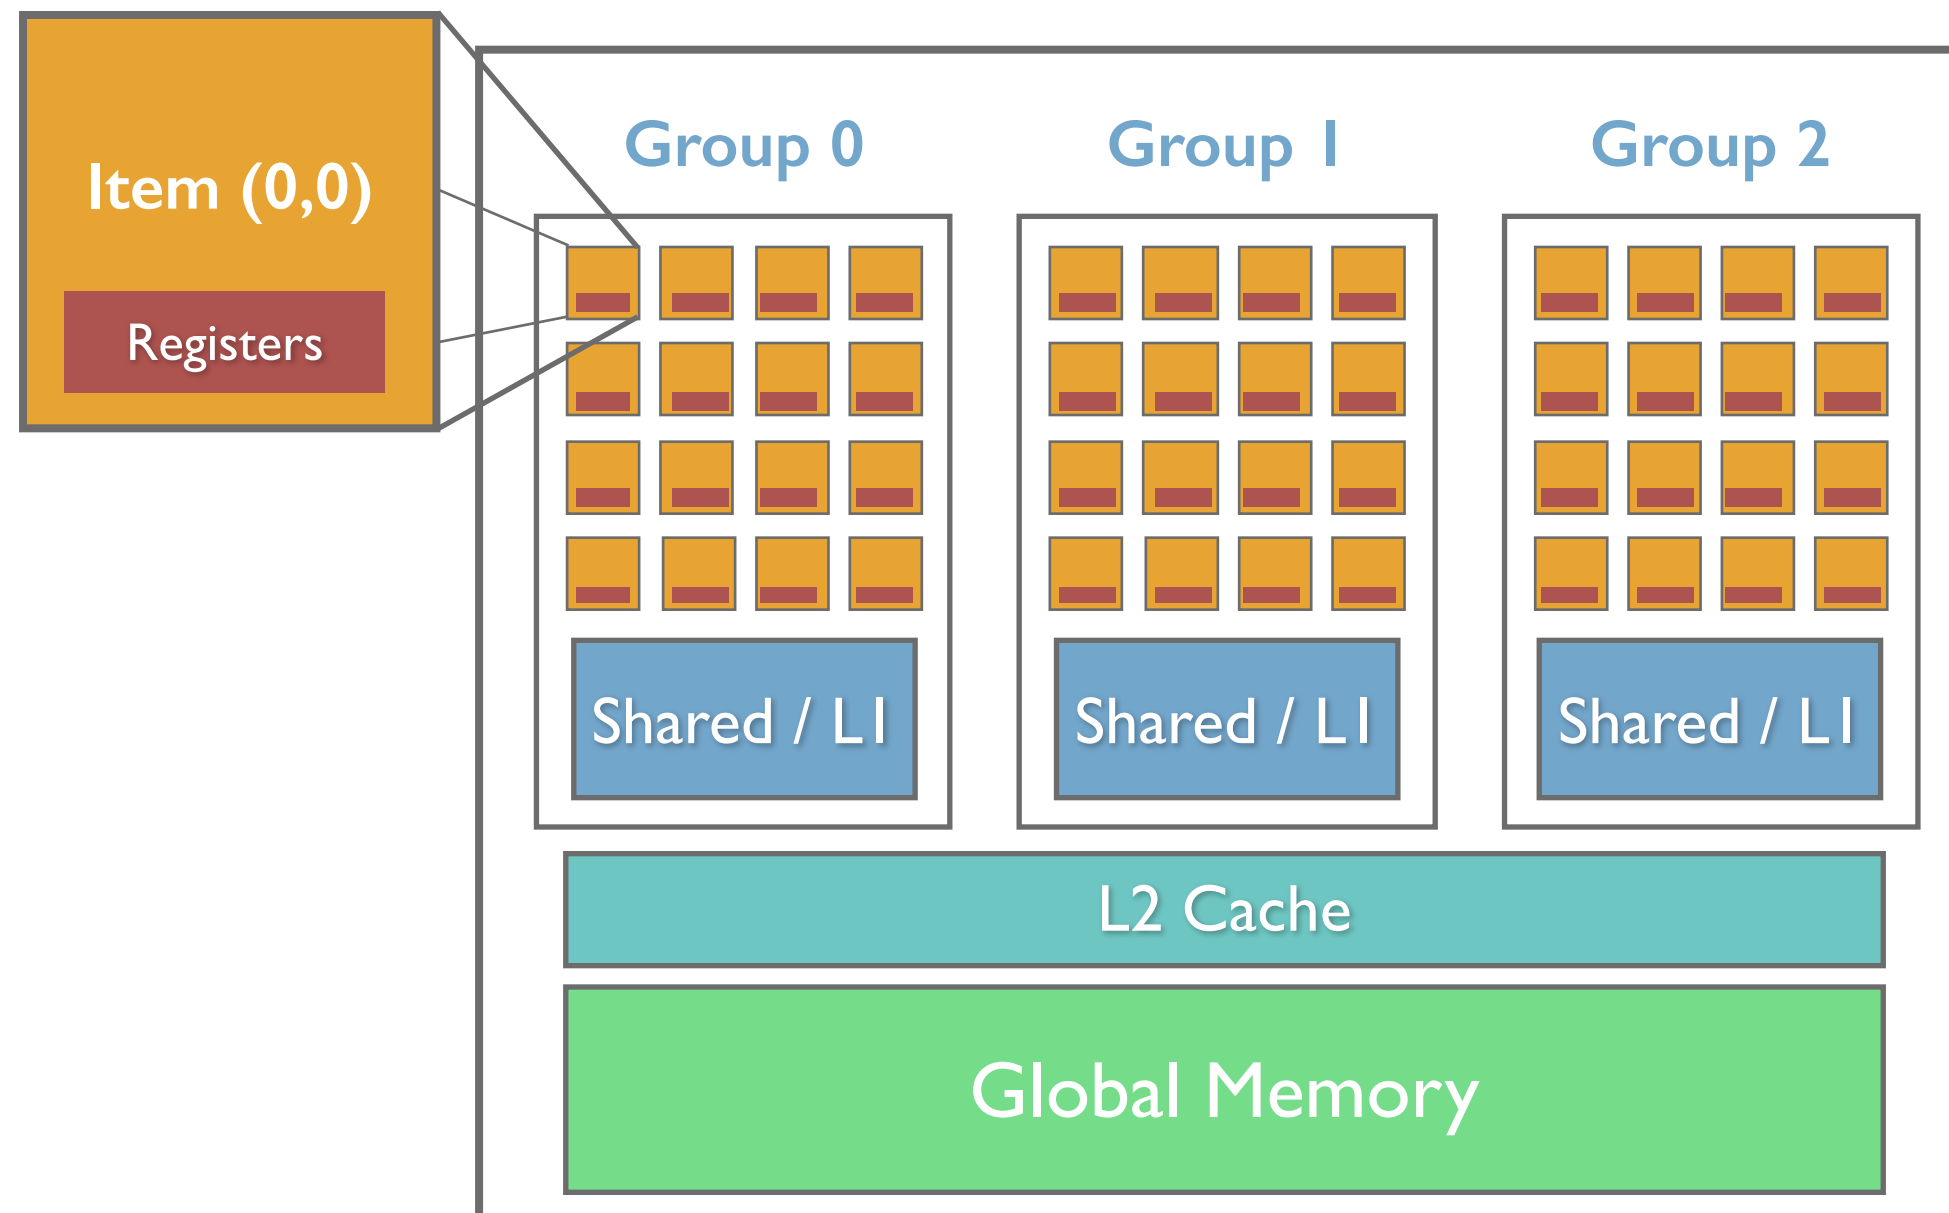
Knights landing die shot  
(removed)

Knights Landing: 72 Silvermont Atom cores  
with 512 bit vector registers. Stacked ram.

# GPU Architecture (abstract)

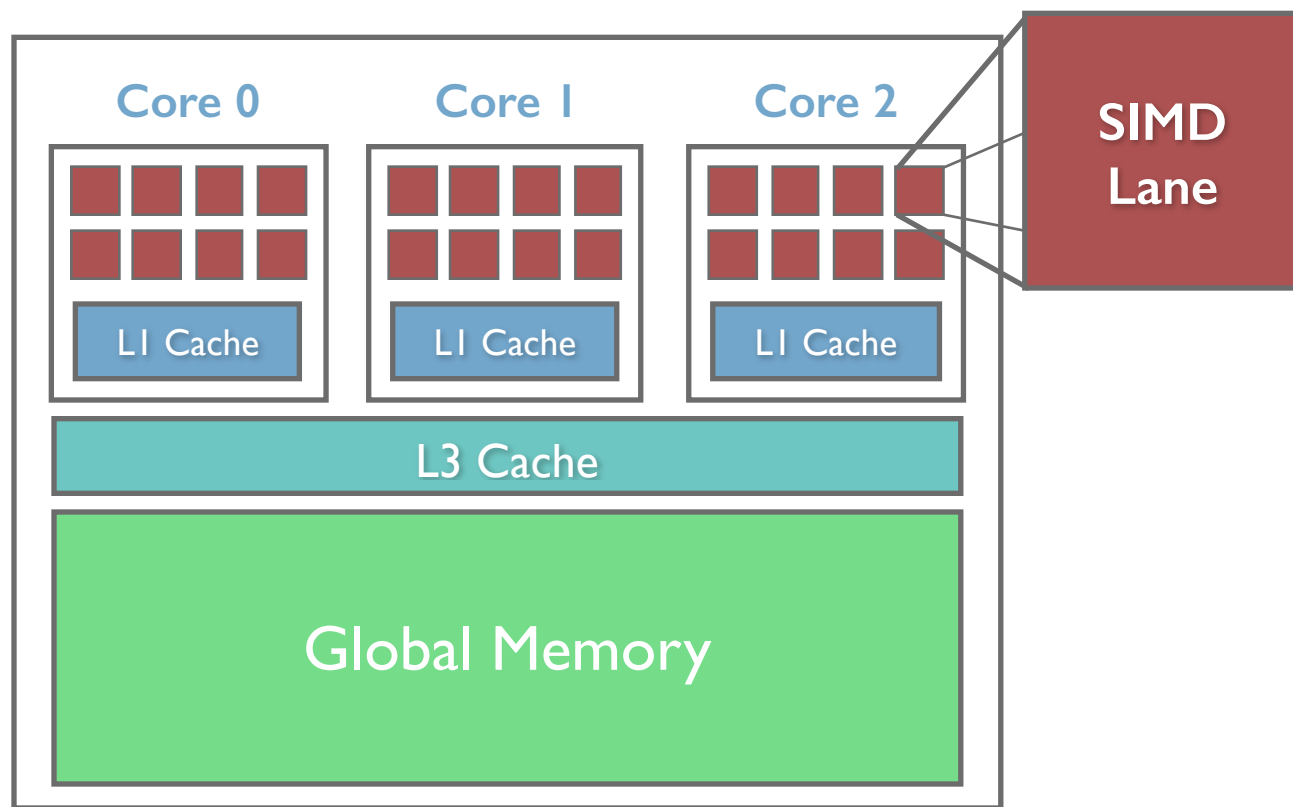
## Parallel model

- Independent work-groups are launched
- Work-groups contain groups of work-items, “parallel” threads.

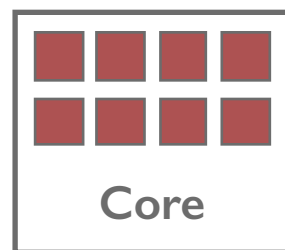
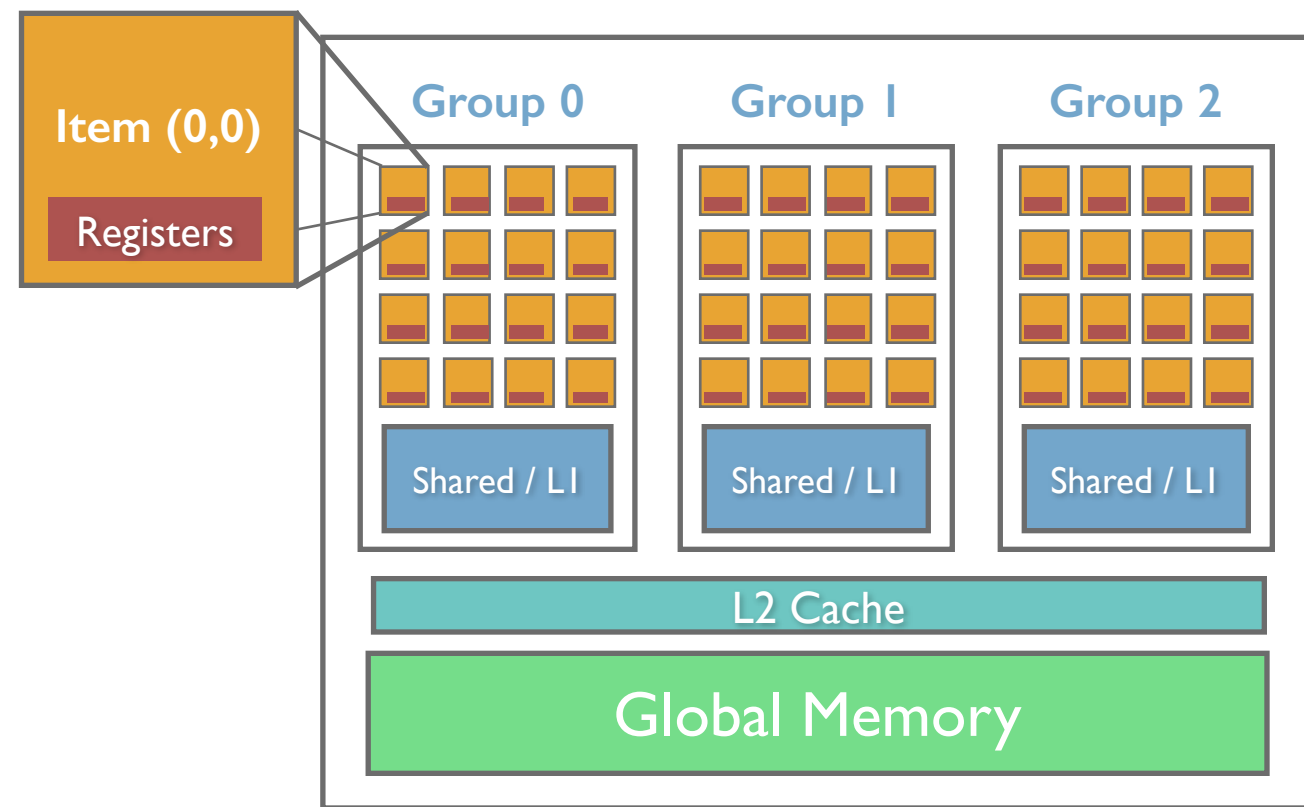


# Parallelization Paradigm

CPU Architecture




GPU Architecture



≈



≈

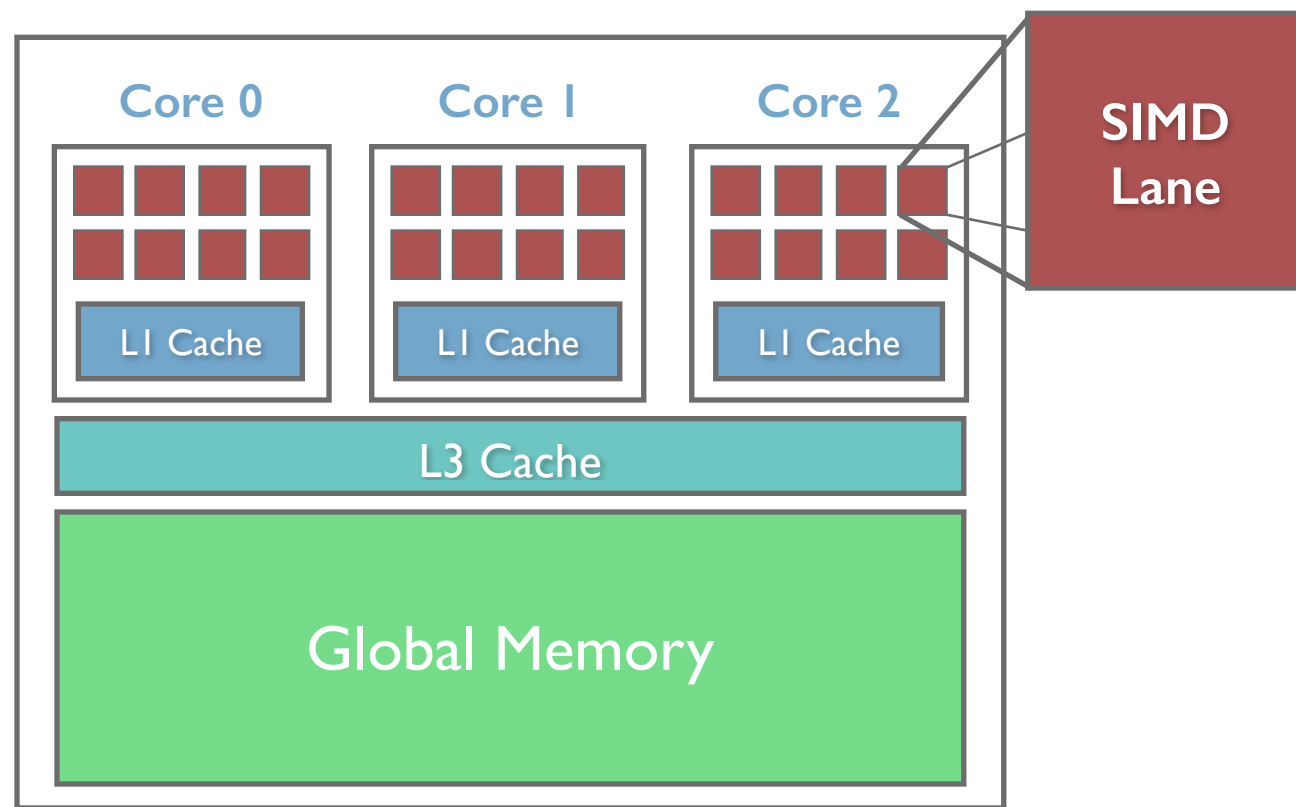


Computational hierarchies are similar

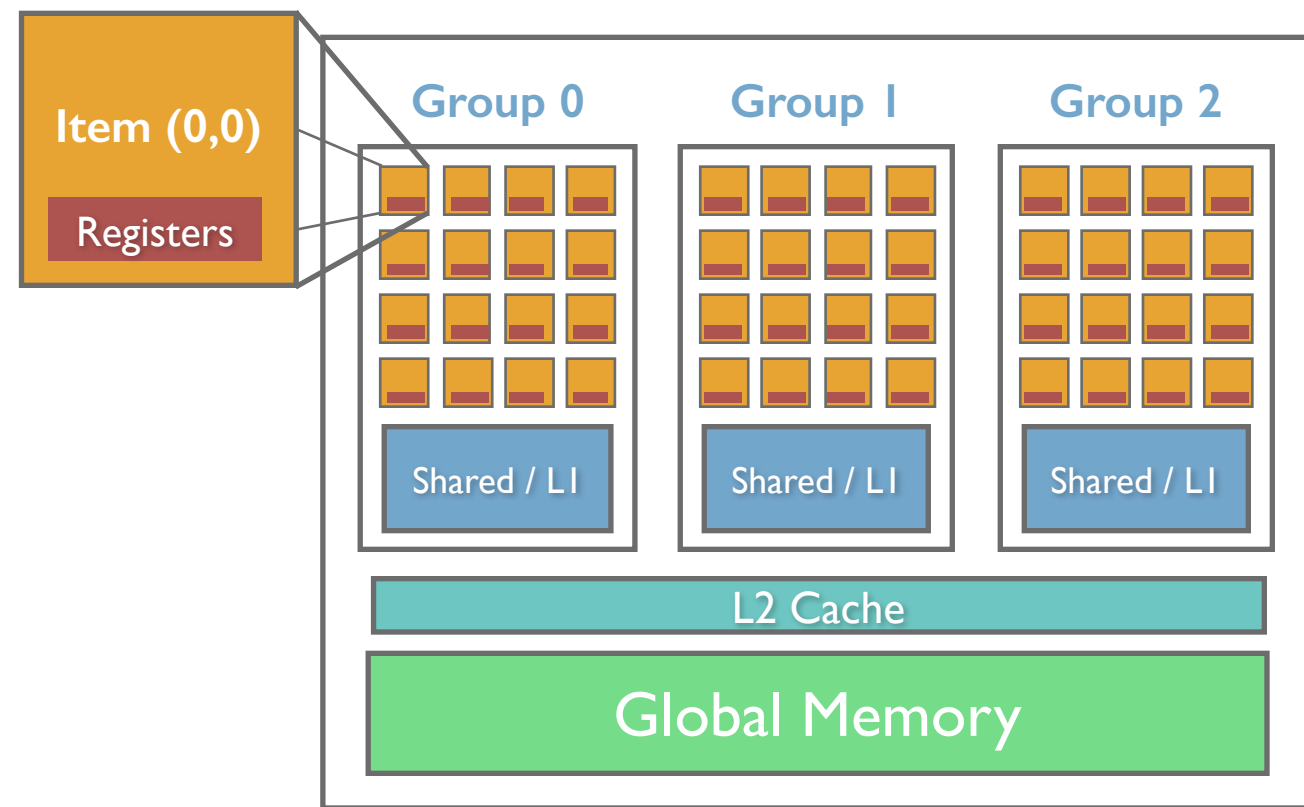


# Parallelization Paradigm

CPU Architecture



GPU Architecture

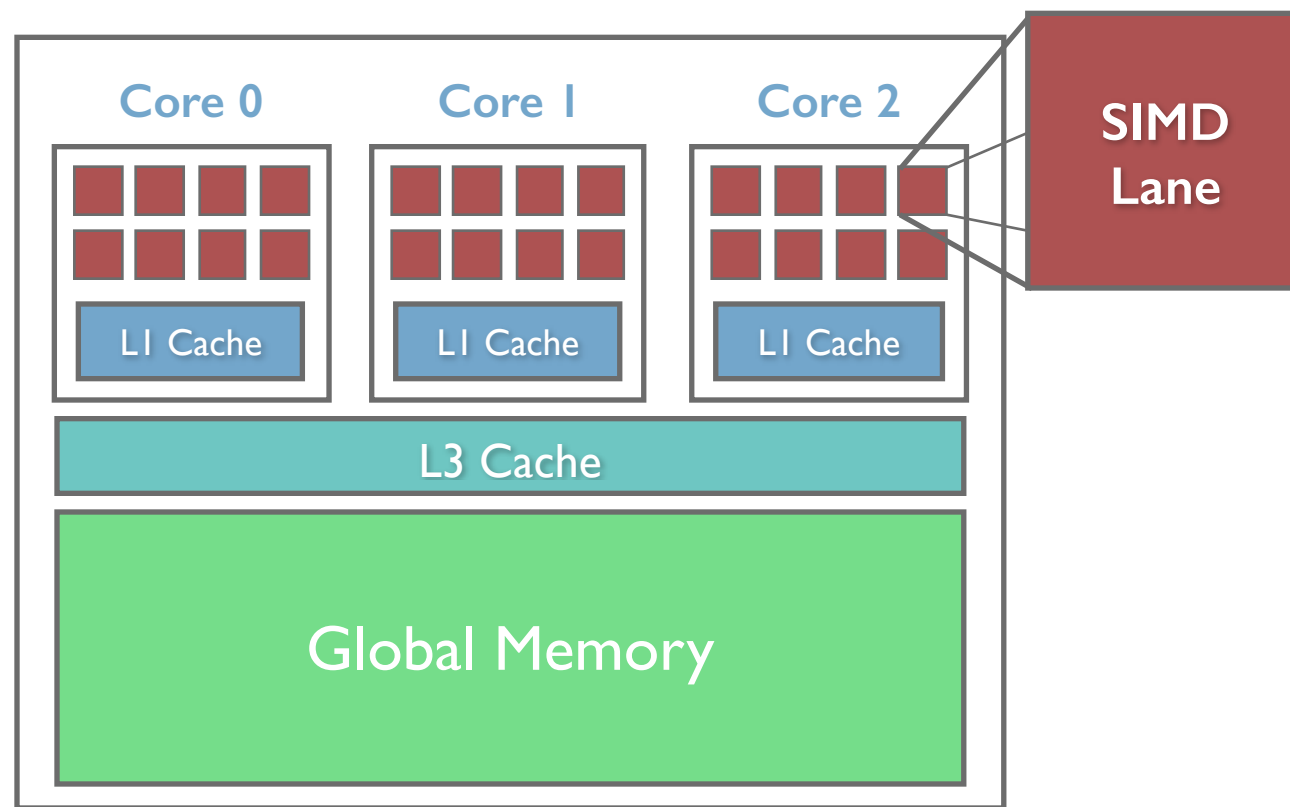


```
void cpuFunction(){  
    #pragma omp parallel for  
    for(int i = 0; i < work; ++i){  
  
        Do [hopefully thread-independent] work  
    }  
}
```

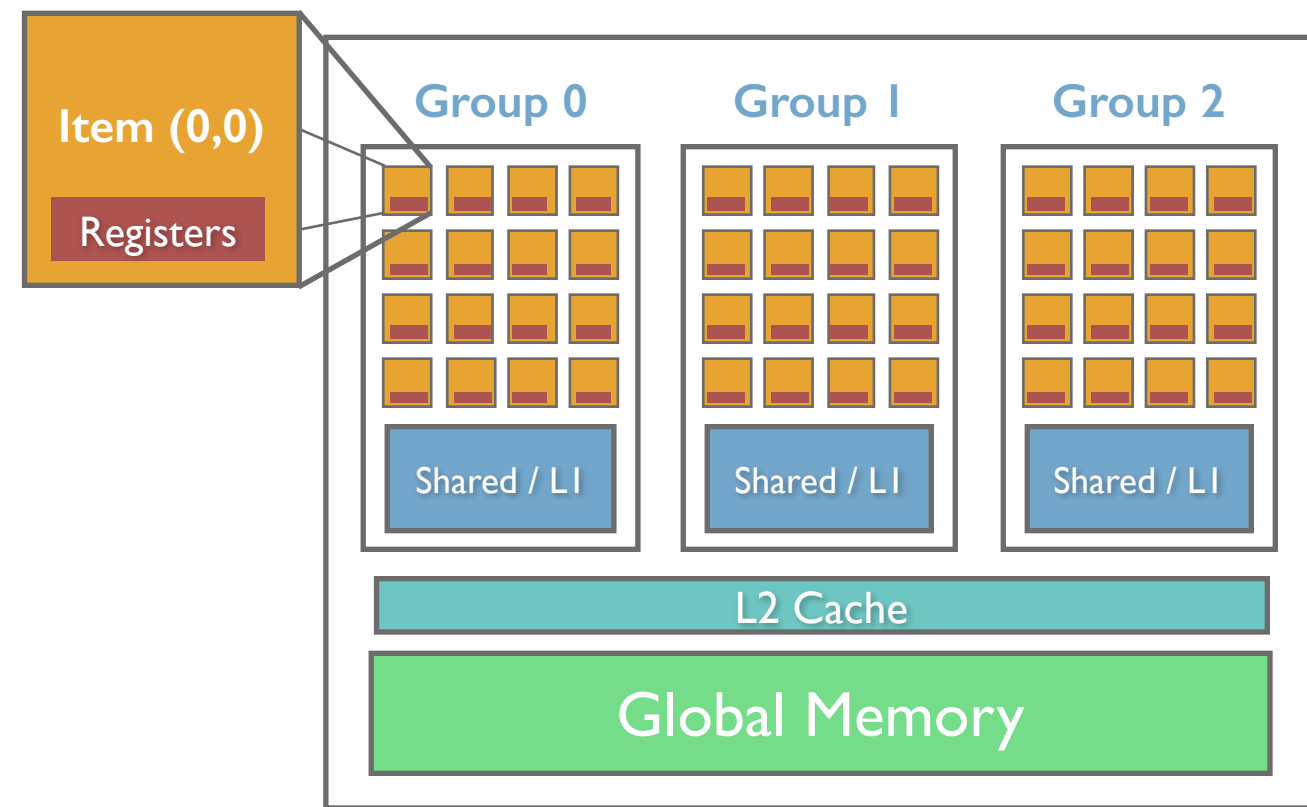
```
__kernel void gpuFunction(){  
    // for each work-group {  
    //     for each work-item in group {  
  
        Do [group-independent] work  
    }  
}
```

# Parallelization Paradigm

CPU Architecture



GPU Architecture

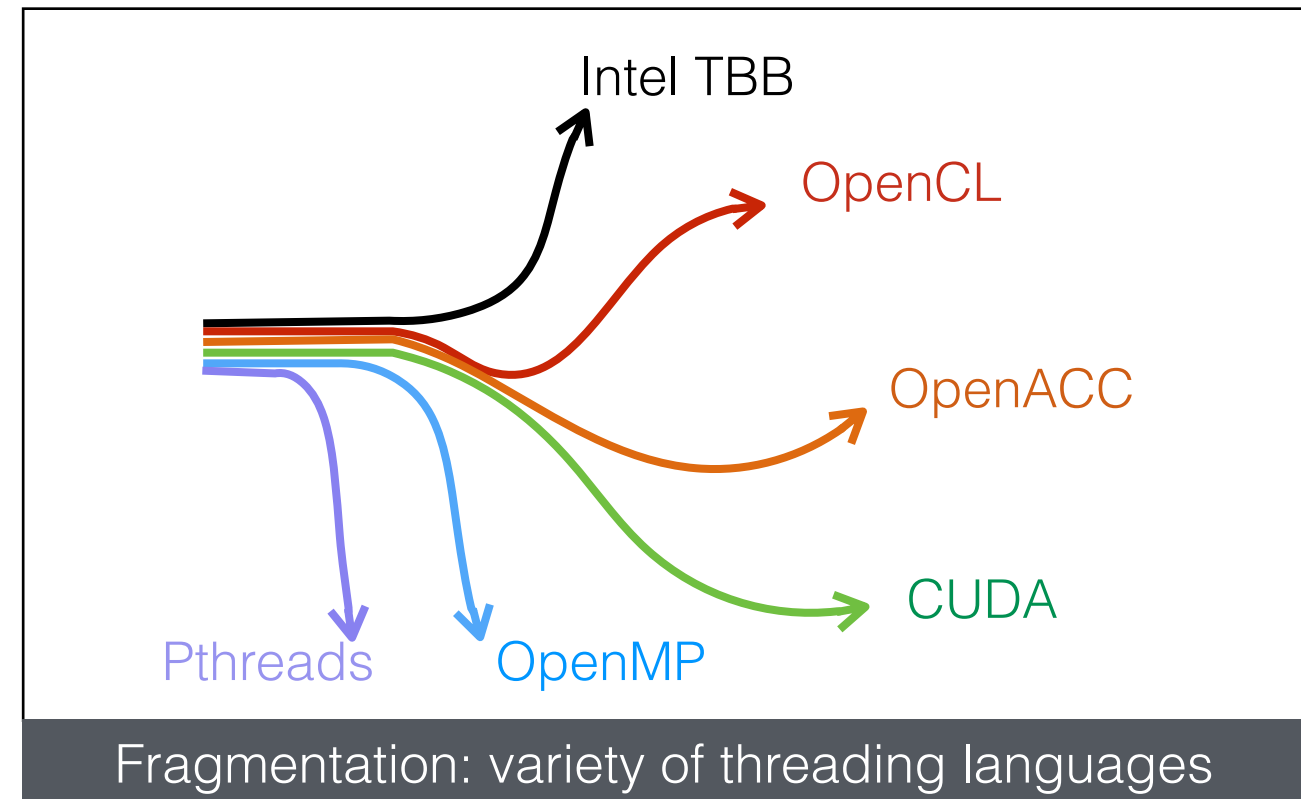
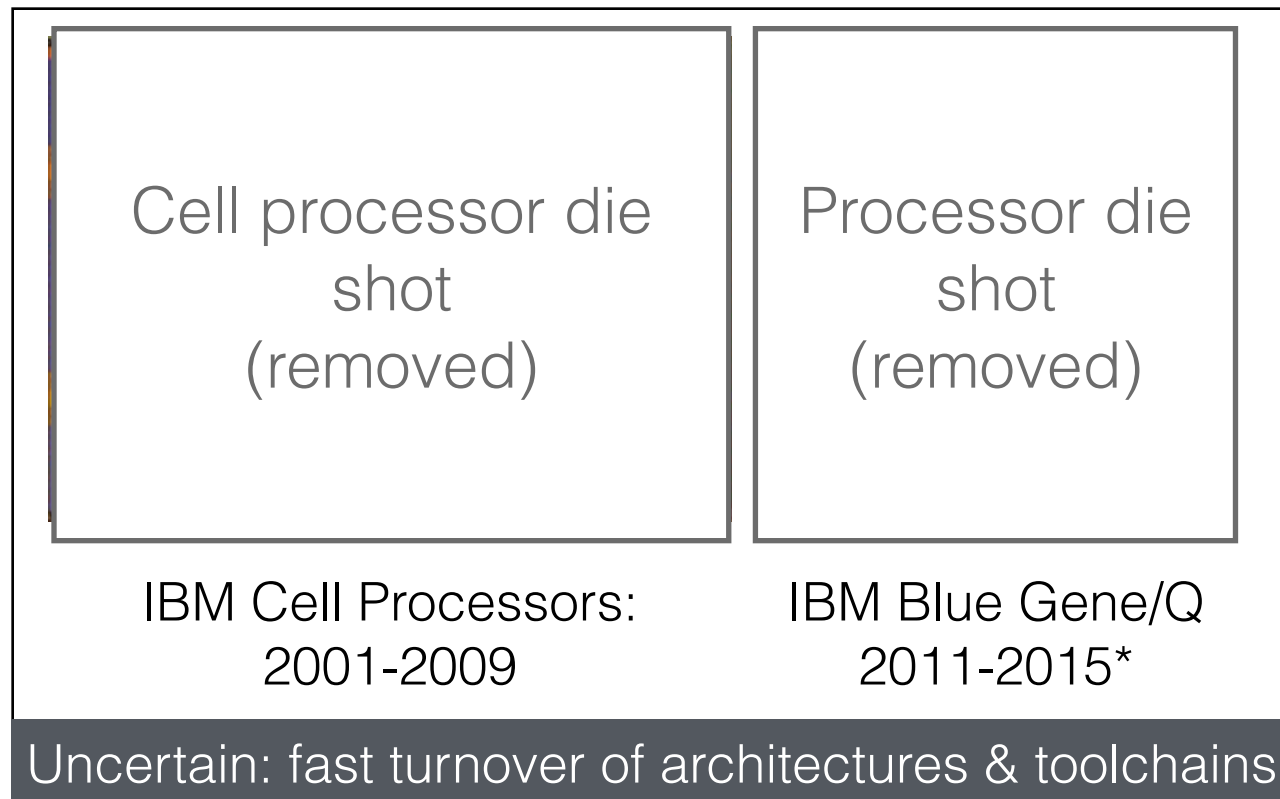


```
void ompFunction(){  
  // for each thread {  
    for(thread's work){  
  
      Do [hopefully thread-independent] work  
  
    }  
  }  
}
```

```
__kernel void gpuFunction(){  
  // for each work-group {  
    // for each work-item in group {  
  
      Do [group-independent] work  
  
    }  
  }  
}
```

# Many-core: challenges

Yet programming massively parallel processors is not mainstream...

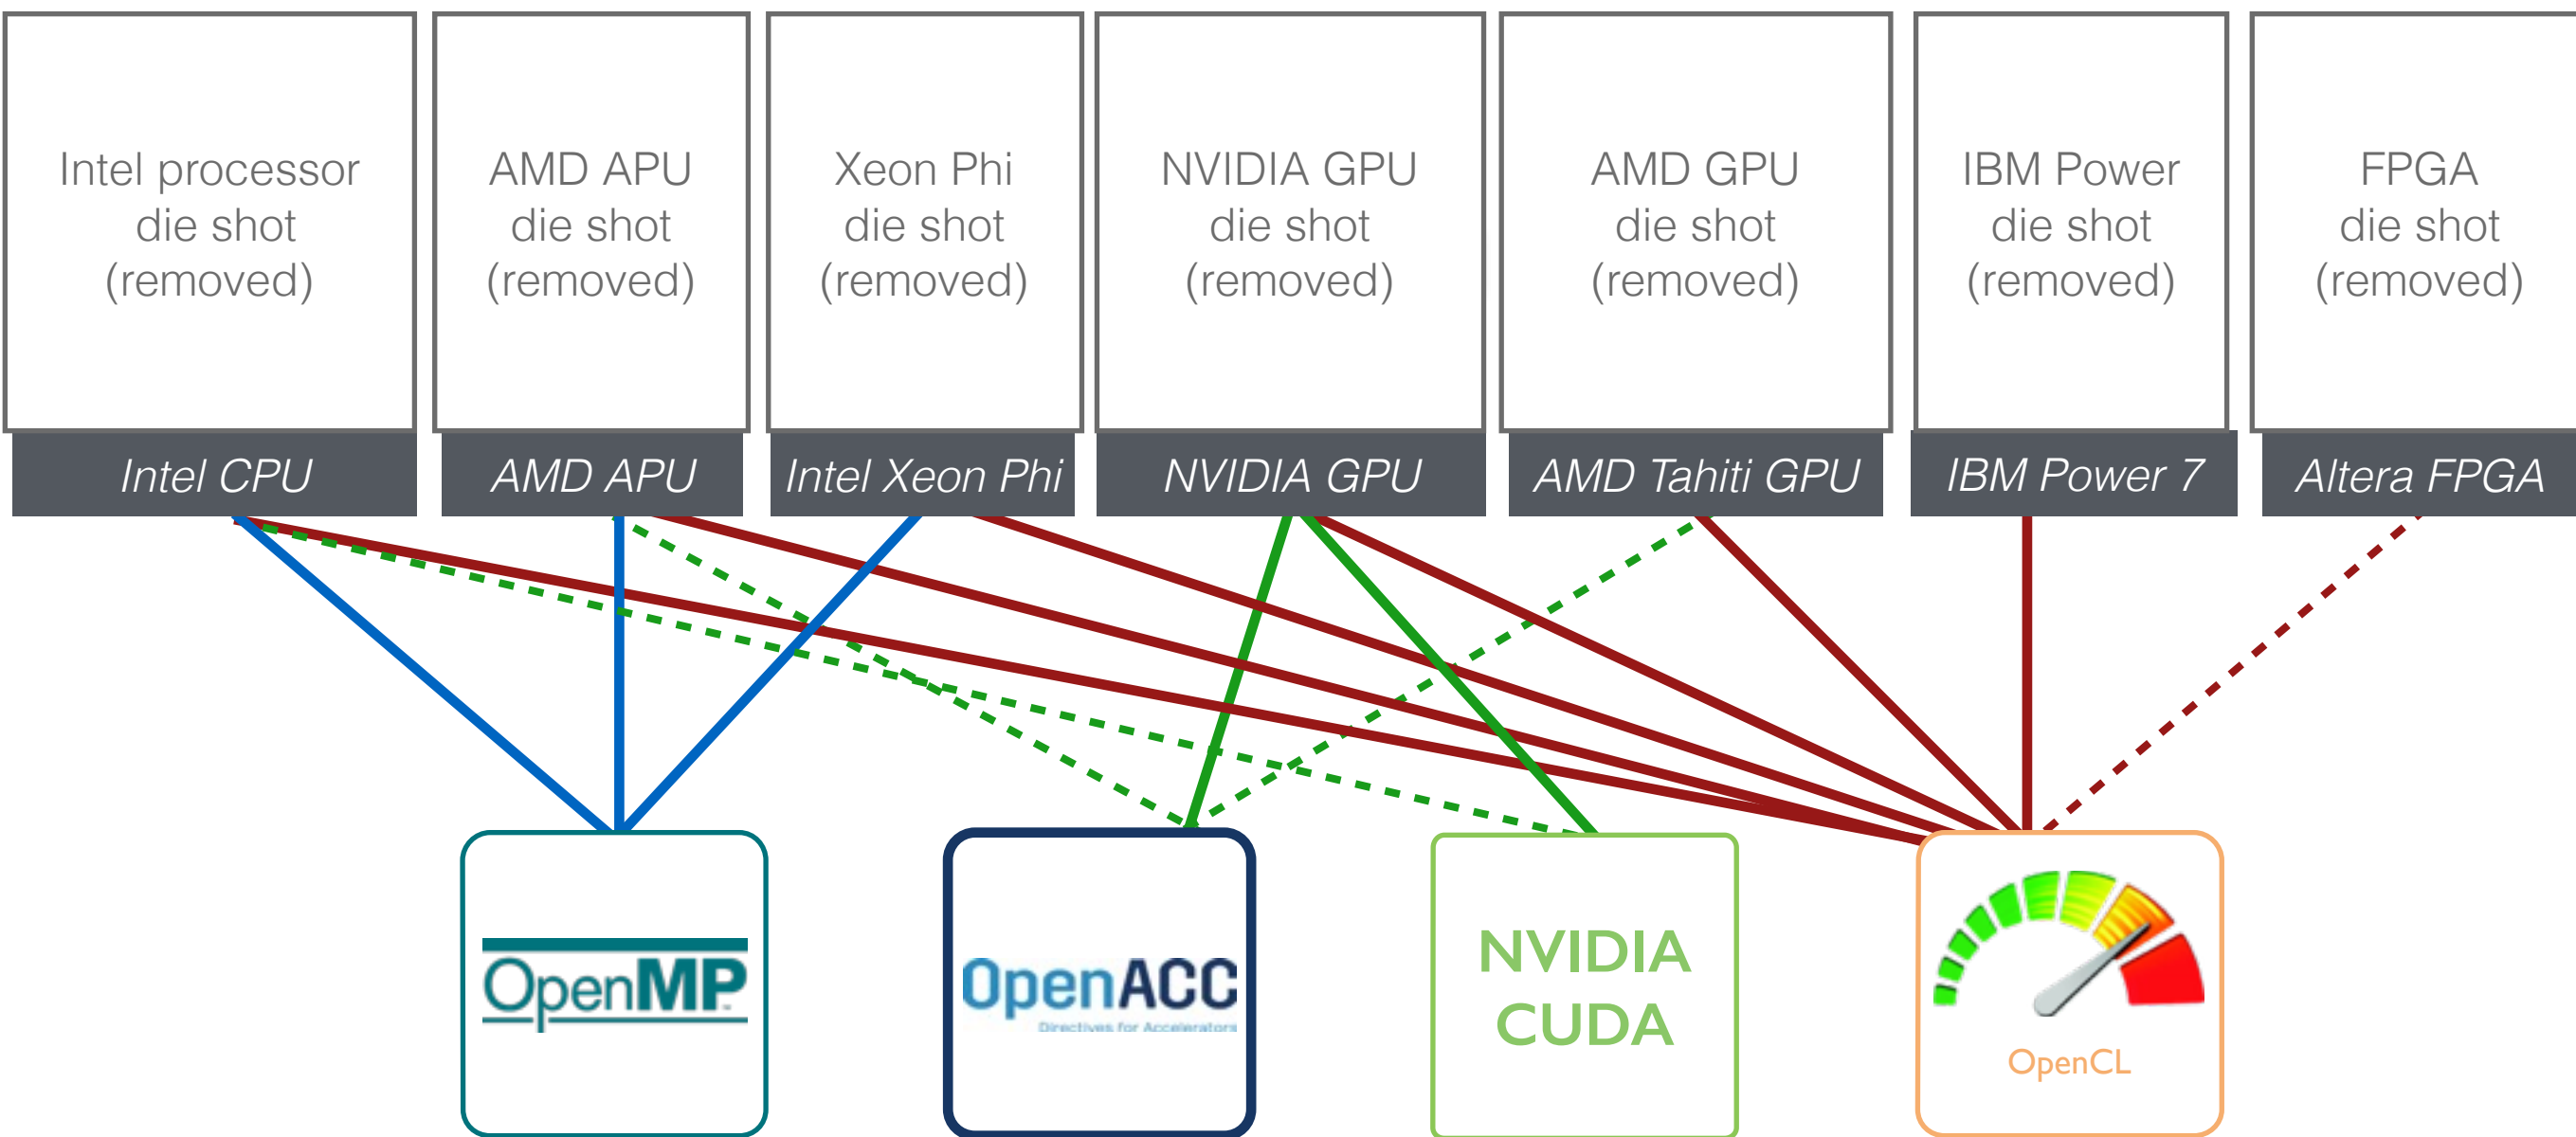


*“In the context of today’s CPU landscape, then, redesigning your application to run multithreaded on a multicore machine is a little like learning to swim by jumping into the deep end”*

The Free Lunch Is Over, Herb Sutter, 2009.

# Many-core: fragmentation

Zoo of competing architectures and programming models (with vendor bias)



Need an efficient, durable, portable, open-source, vendor-independent approach for many-core programming



# Motivation

## Uncertainty

- Code life cycle measured in decades
- Architecture & API life cycles measured in Moore doubling periods
- Example: IBM Cell processor, IBM Blue Gene Q

## Portability

- CUDA, OpenCL, OpenMP, OpenACC, Intel TBB... are not code compatible
- Not all APIs are installed on any given system

## Performance

- Logically similar kernels differ in performance (GCC & ICPC, **OpenCL & CUDA**)
- Naively porting OpenMP to CUDA or OpenCL will likely yield low performance

## Programmability

- Expose parallel paradigm ... without introducing an exotic programming model

# Portability Approaches

## Directive approach

- Use of optional [`#pragma`]'s to give compiler transformation hints
- Aims for portability, **performance** and programmability

## Source-to-source approach

- Compiler tools can be used to translate across specifications/languages
- **Performance** is not always **portable**
- **Maintenance** of original and translated codes

## Wrapper approach

- Create a tailored library with **optimized** functions
- **Restricted** to a pre-canned set of operations
- Flexibility comes from functors/lambda's at **compile-time**

# Portability Approaches

## Directive approach

- Use of optional `[#pragma]`'s to give compiler transformation hints
- Aims for portability, **performance** and programmability



- Introduced for accelerator support through directives (2012)
- There are compilers which support the 1.0 specifications
- OpenACC 2.0 introduces support for inlined functions



- OpenMP has been around for a while (1997)
- OpenMP 4.0 specifications (2013) includes accelerator support
- Few compilers (ROSE) support parts of the 4.0 specifications



```
#pragma omp target teams distribute parallel for
for(int i = 0; i < N; ++i){
    y[i] = a*x[i] + y[i];
}
```

# Portability Approaches

## Directive approach

- Not centralized anymore due to the offload model
- OpenACC and OpenMP begin to resemble an API rather than code decorations

OpenACC  
Directives for Accelerators

```
double a[100];  
#pragma acc enter data copyin(a)  
// OpenACC code  
#pragma acc exit data copyout(a)
```

OpenACC  
Directives for Accelerators

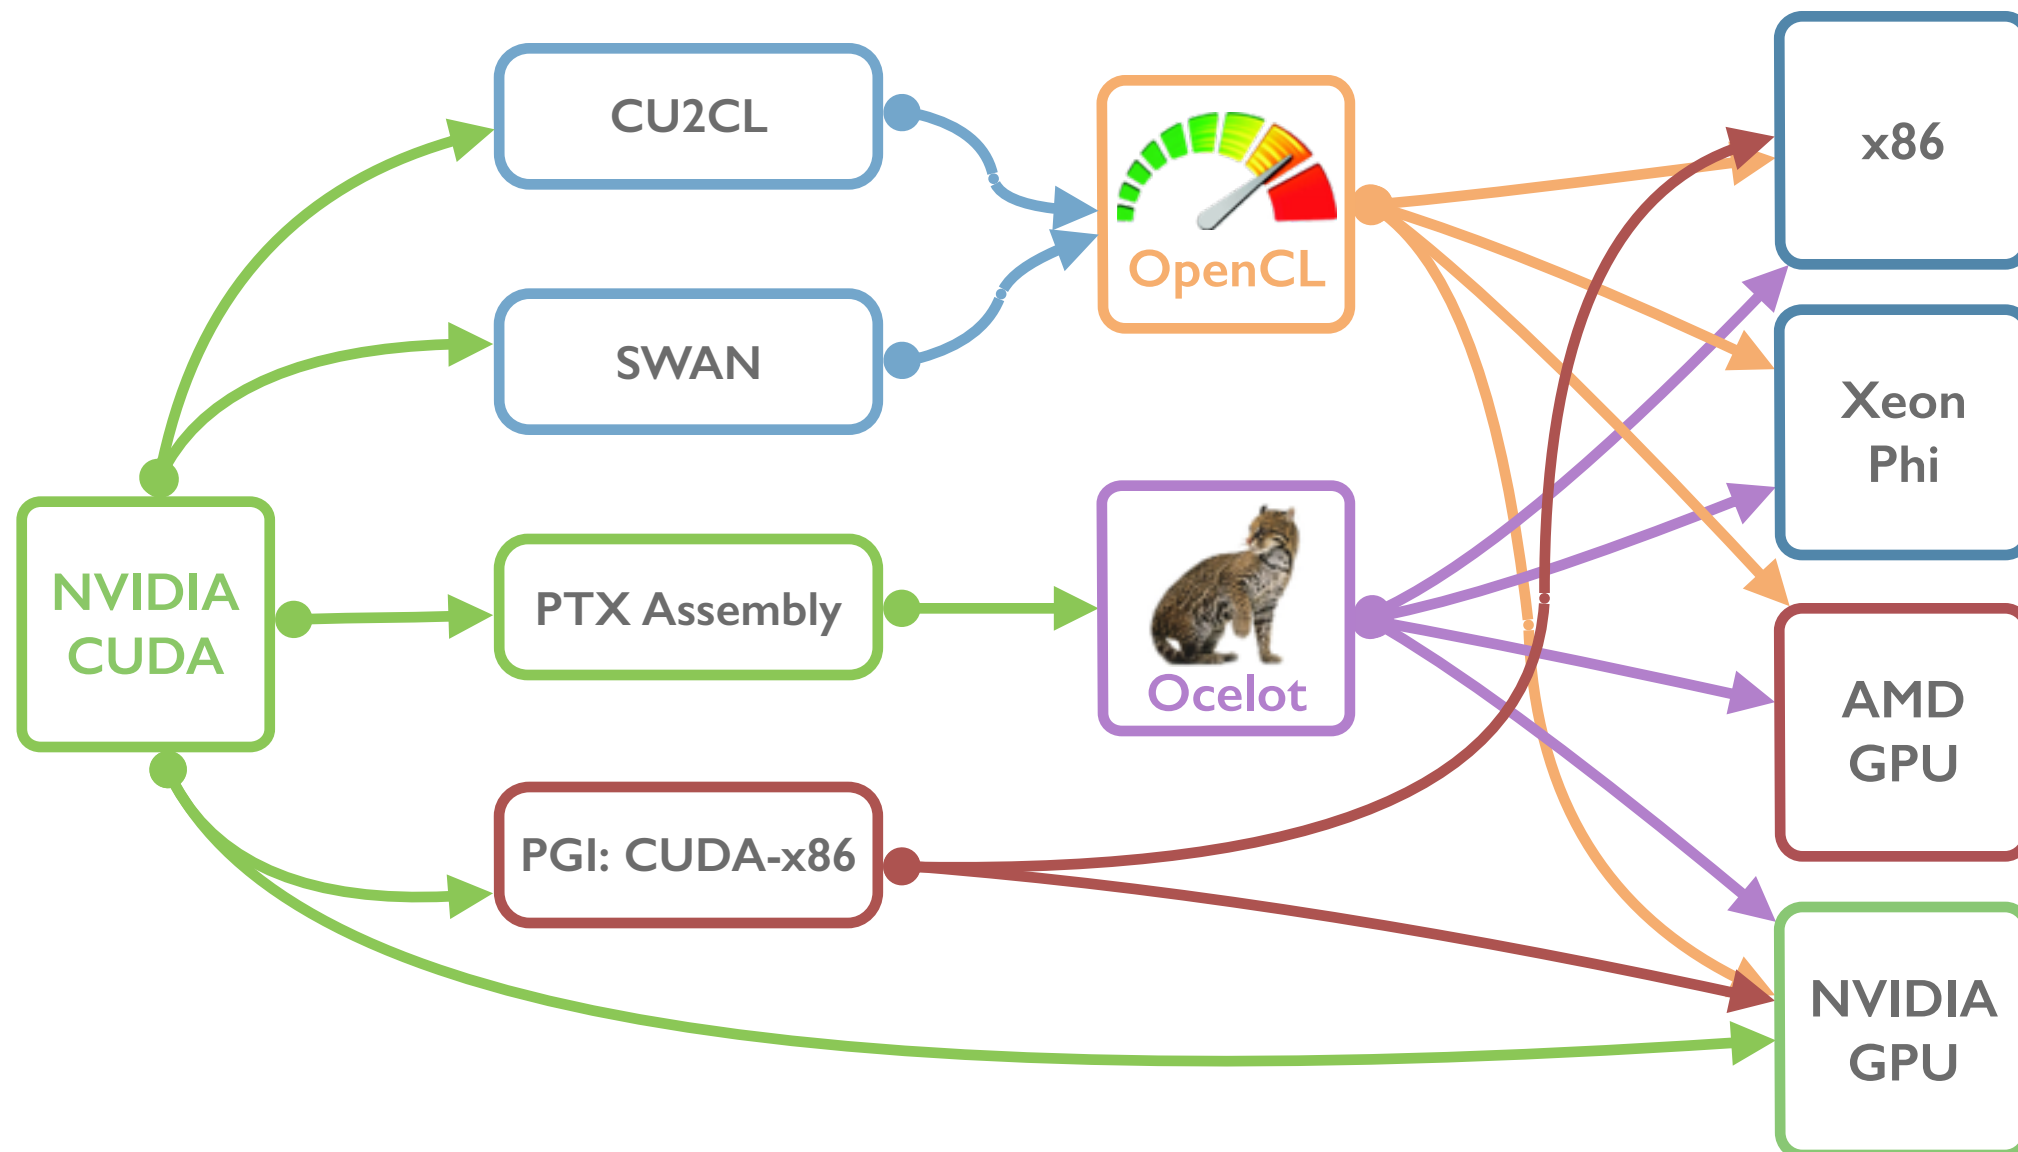
```
class Matrix {  
    double *v;  
    int len  
    Matrix(int n) {  
        len = n;  
        v = new double[len];  
        #pragma acc enter data create(v[0:len])  
    }  
    ~Matrix() {  
        #pragma acc exit data delete(v[0:len])  
        delete[] v;  
    }  
};
```



# Portability Approaches

## Source-to-source approach

- CU2CL and SWAN have limited CUDA support (3.2 and 2.0 respectively)
- GPU Ocelot supports PTX from CUDA 4.2 (5.0 partially)
- PGI: CUDA-x86 appears to have been put in hiatus since 2011



GPU Ocelot seems to be the most active of these projects

# Portability Approaches

## Wrapper approach

- Create a tailored library with **optimized** functions
- **Restricted** to a set of operations with flexibility from functors/lambda



- C++ library masking OpenMP, Intel's TBB and CUDA for x86 processors and NVIDIA GPUs
- **Vector library**, such as the standard template library (STL)



- Kokkos is from Sandia National Laboratories
- C++ **vector library** with linear algebra routines
- Uses OpenMP and CUDA for x86 and NVIDIA GPU support



- C++ template library
- Uses **code skeletons** for map, reduce, scan, mapreduce, ...
- Uses OpenMP, OpenCL and CUDA as backends

# Portability Approaches

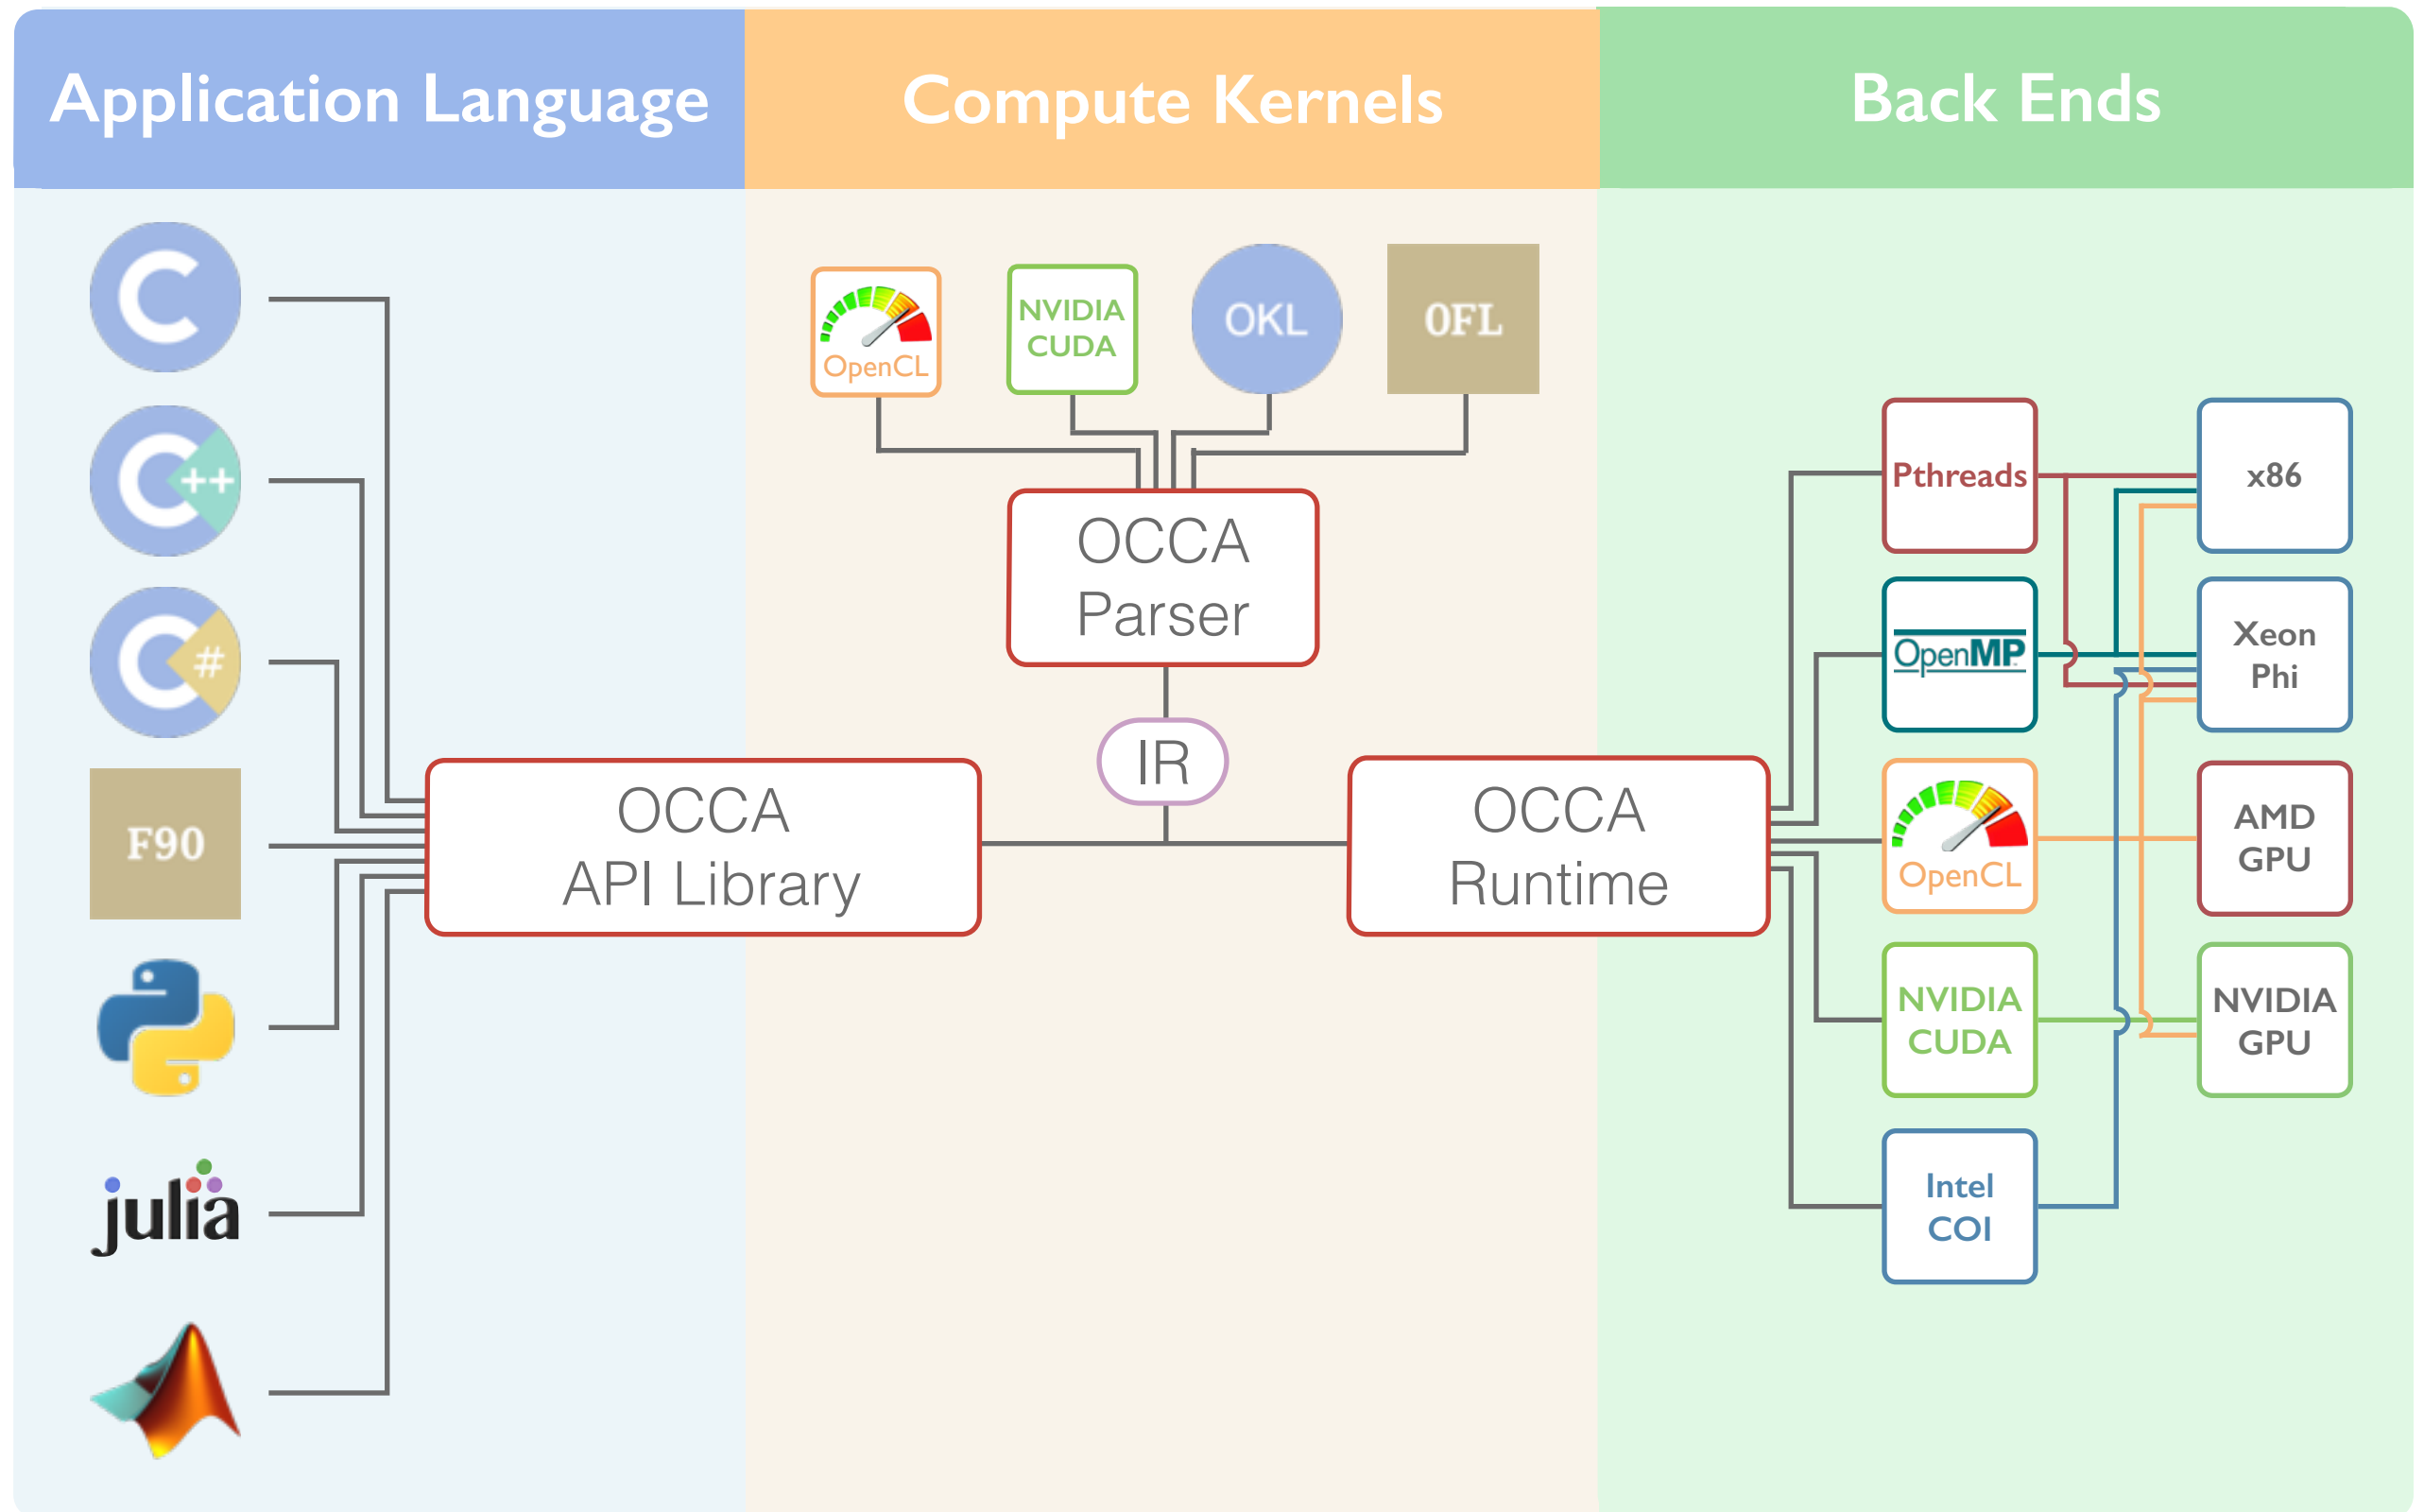
Numerous approaches to portability

API	Type	Front-ends	Kernel	Back-ends	DSL
Kokkos	ND arrays	C++	Custom	CUDA & OpenMP	
VexCL	Vector class	C++	-	CUDA & OpenCL	
RAJA	Library	C++	C++ Lambdas	CUDA, OpenMP, OpenACC	
→ OCCA2	API, Source-to-source, Kernel Languages	C,C++,C#, F90, Python,MATLAB, Julia	OpenCL, CUDA,& custom unified kernel language	CUDA, OpenCL, pThreads,OpenMP, Intel COI	Compiler
CU2CL *	Source-to-source	App	CUDA	OpenCL	
Insieme	Source-to-source compiler	C	OpenMP,Cilk, MPI, OpenCL	OpenCL,MPI, Insieme IR runtime	
Trellis	Directives	C/C++	#pragma trellis	OpenMP, OpenACC, CUDA	
OmpSs	Directives + kernels	C,C++	Hybrid OpenMP, OpenCL, CUDA	OpenMP, OpenCL, CUDA	
Ocelot	PTX Translator	CUDA	CUDA	OpenCL	

OCCA emphasis: lightweight and extensible.

\* Wu Feng et al @ VT !

# OCCA: accessible portability





# OCCA Application Programming Interface (API)

`occa::device` (C++ Class)

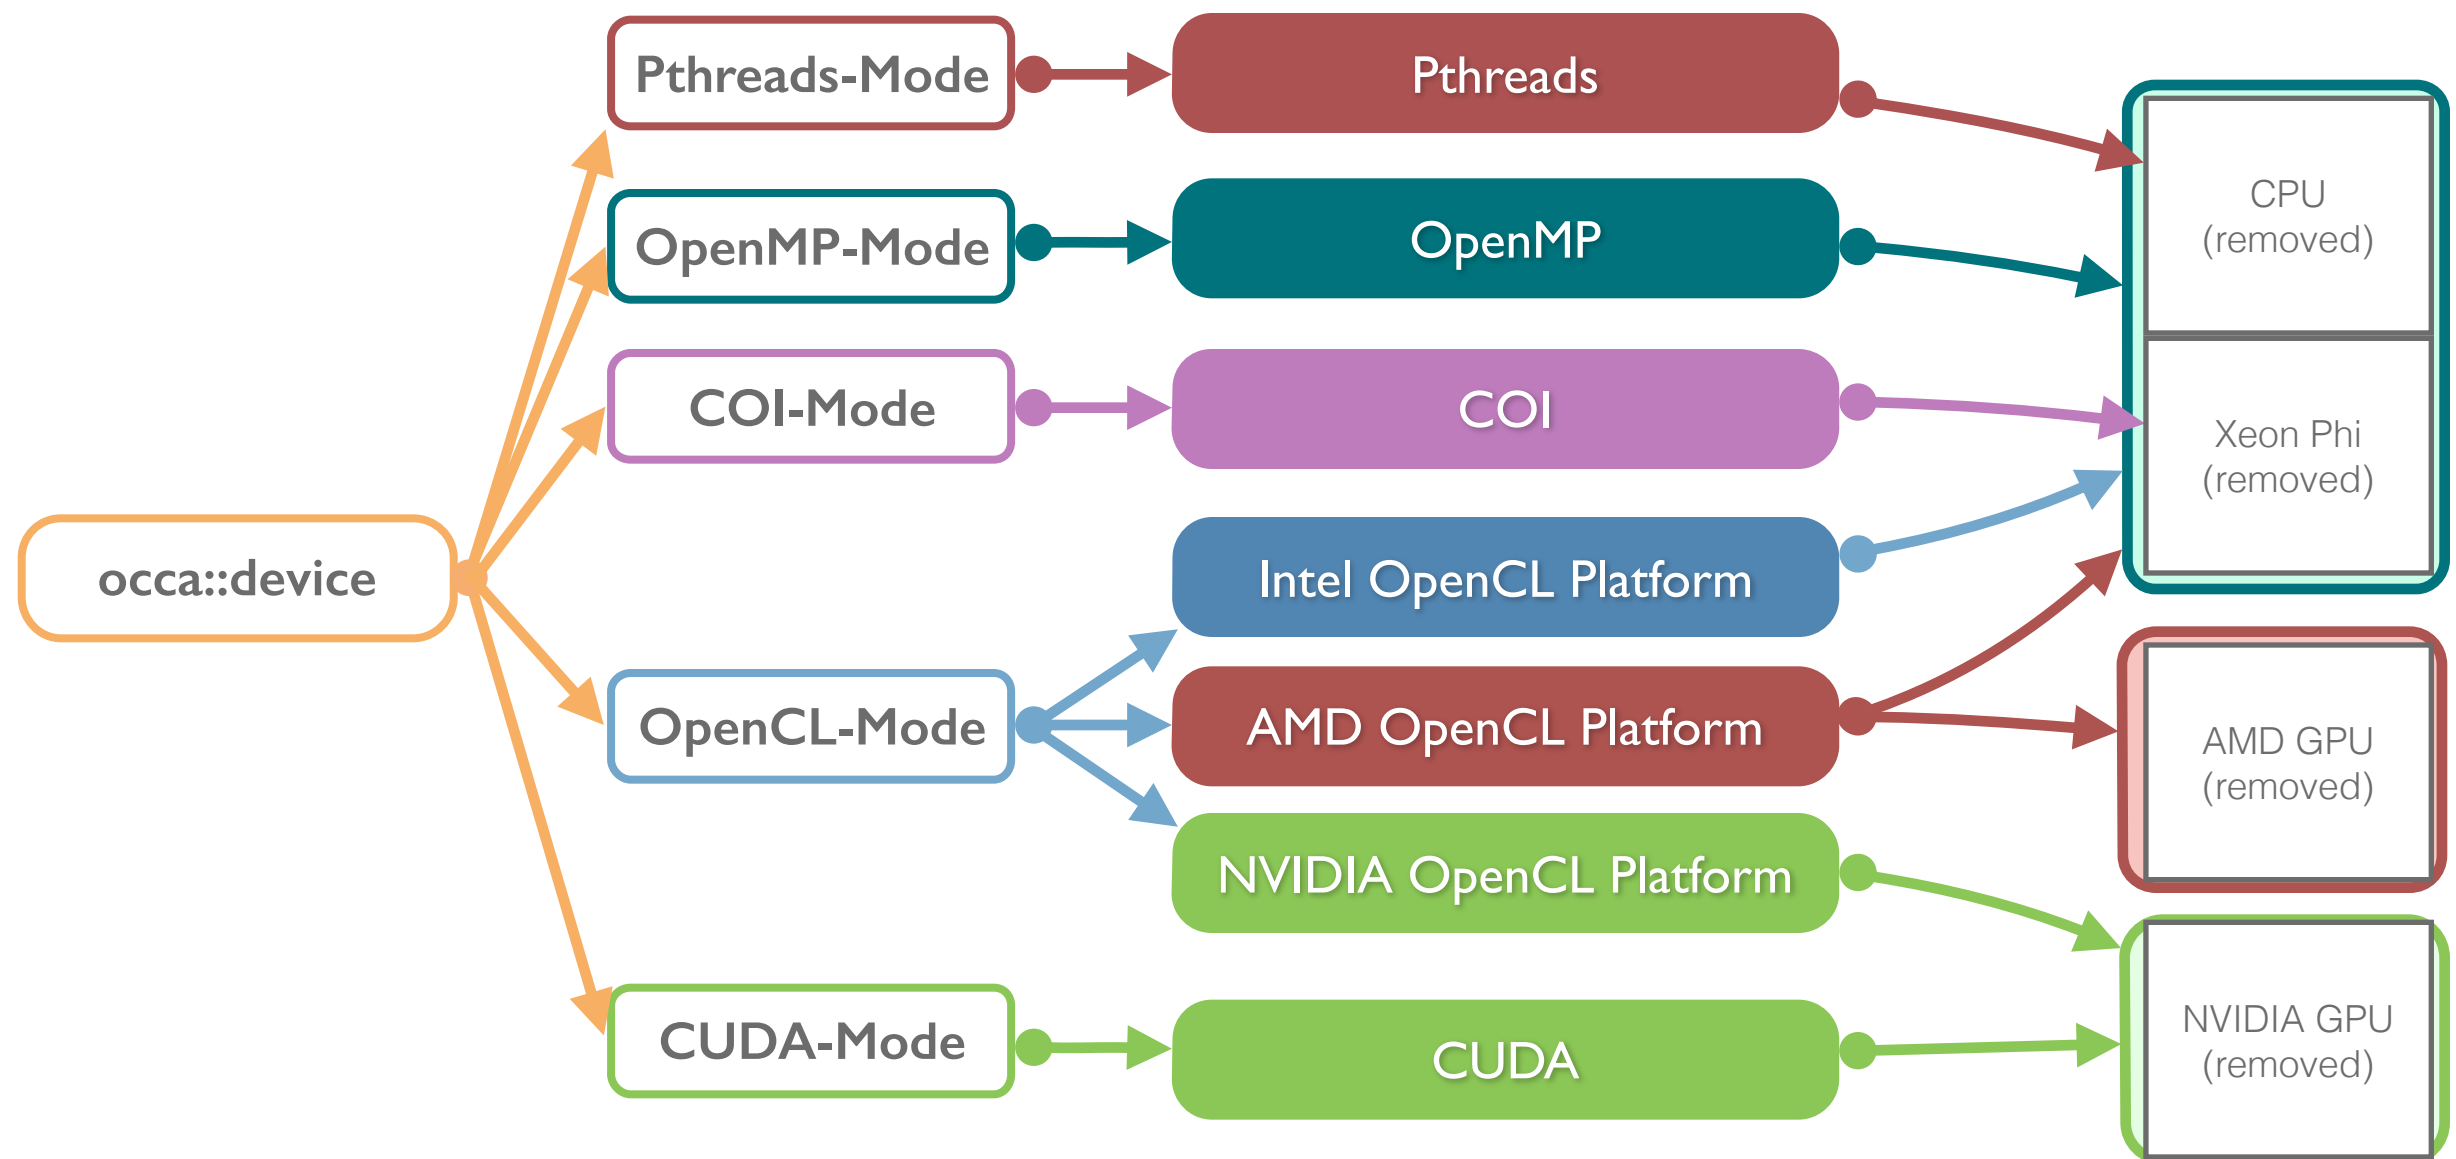
`occa::memory` (C++ Class)

`occa::kernel` (C++ Class)

# OCCA Application Programming Interface (API)

`occa::device` (C++ Class)

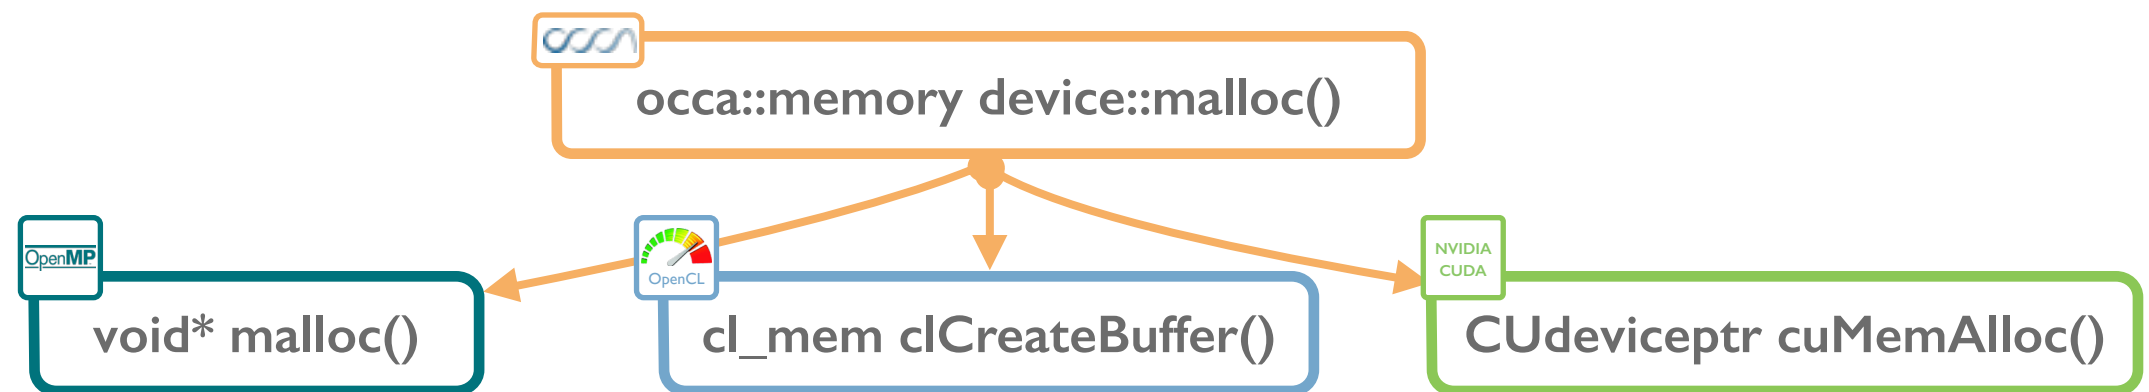
- Choose between using the CPU or available accelerators
- In charge of allocating memory and compiling kernels



# OCCA Application Programming Interface (API)

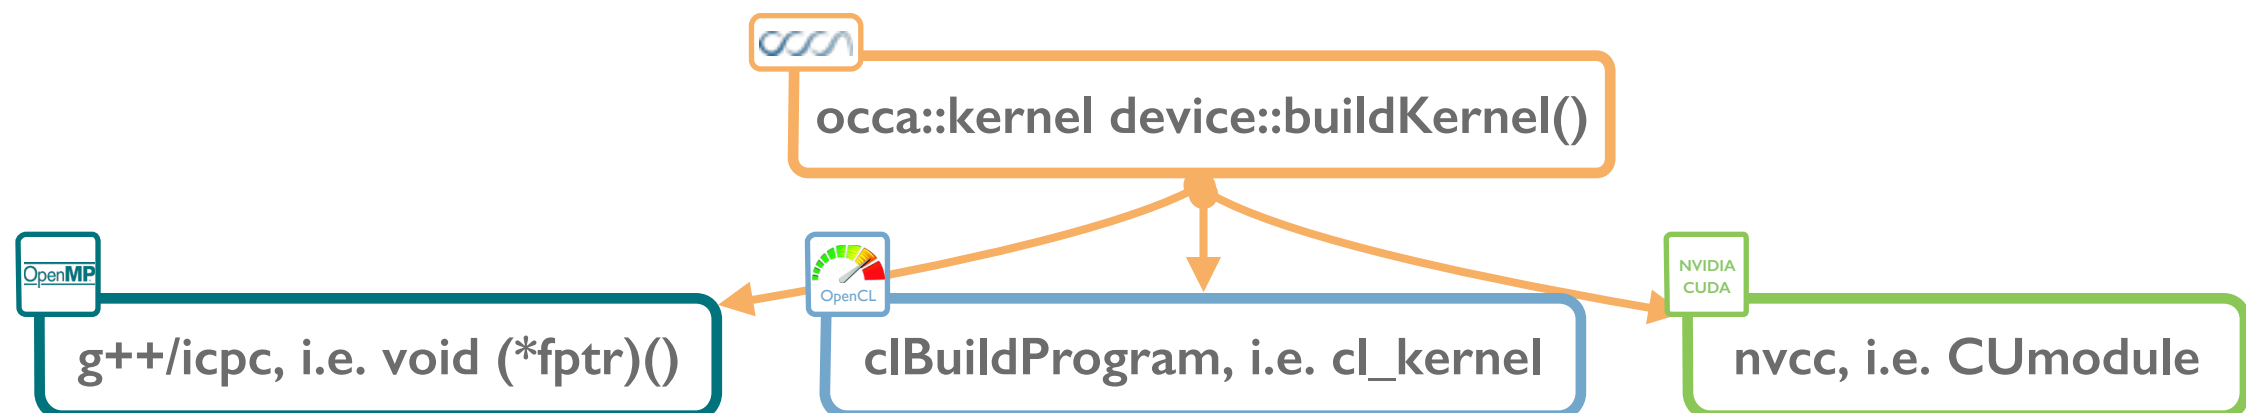
## occa::memory (C++ Class)

- Abstracts the memory handles found in each language
- Asynchronous memory transfers are supported



## occa::kernel (C++ Class)

- Uses run-time compilation
- Kernel binaries are cached to prevent re-compiling



# OKL: OCCA Kernel Language

## Description

- Minimal extensions to C, familiar for regular programmers
- Explicit loops expose parallelism for modern multicore CPUs and accelerators
- Parallel loops are explicit through the fourth for-loop **inner** and **outer** labels

```
kernel void kernelName(...){  
    ...
```

```
    for(int groupZ = 0; groupZ < zGroups; ++groupZ; outer2){  
        for(int groupY = 0; groupY < yGroups; ++groupY; outer1){  
            for(int groupX = 0; groupX < xGroups; ++groupX; outer0){
```

```
                for(int itemZ = 0; itemZ < zItems; ++itemZ; inner2){  
                    for(int itemY = 0; itemY < yItems; ++itemY; inner1){  
                        for(int itemX = 0; itemX < xItems; ++itemX; inner0){  
                            // GPU Kernel Scope
```

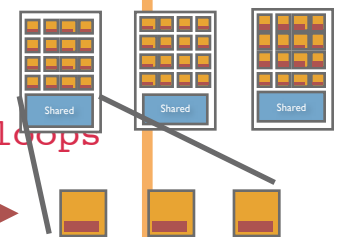
```
                        }  
                    }  
                }
```

```
            }  
        }  
    }
```

```
    ...  
}
```

// Work-group implicit loops

// Work-item implicit loops



NVIDIA  
CUDA

```
dim3 blockDim(xGroups,yGroups,zGroups);  
dim3 threadDim(xItems,yItems,zItems);  
kernelName<<< blockDim , threadDim >>>(...);
```

OKL



# OKL: OCCA Kernel Language

## Outer-loops

- **Outer-loops** are synonymous with CUDA and OpenCL **kernels**
- Extension: allow for multiple outer-loops per kernel

```
kernel void kernelName(...){  
    ...  
  
    for(outer){  
        for(inner){  
        }  
    }  
  
    ...  
}
```

OKL

# OKL: OCCA Kernel Language

## Outer-loops

- **Outer-loops** are synonymous with CUDA and OpenCL **kernels**
- Extension: allow for multiple outer-loops per kernel

```
kernel void kernelName(...){  
    for(outer){  
        for(inner){  
            }  
        }  
  
    for(outer){  
        for(inner){  
            }  
        }  
  
    for(outer){  
        for(inner){  
            }  
        }  
    }  
}
```

OKL

# OKL: OCCA Kernel Language

## Outer-loops

- **Outer-loops** are synonymous with CUDA and OpenCL **kernels**
- Extension: allow for multiple outer-loops per kernel

```
kernel void kernelName(...) {  
  
    if(expr) {  
        for(outer) {  
            for(inner) {  
                }  
            }  
        }  
    else {  
        for(outer) {  
            for(inner) {  
                }  
            }  
        }  
    }  
  
    while(expr) {  
        for(outer) {  
            for(inner) {  
                }  
            }  
        }  
    }  
}
```

OKL

# OKL: OCCA Kernel Language

## Shared memory

```
for(int groupX = 0; groupX < xGroups; ++groupX; outer0){ // Work-group implicit loops
    shared int sharedVar[16];

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        sharedVar[itemX] = itemX;
    }

    // Auto-insert [barrier(localMemFence);]

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        int i = (sharedVar[itemX] + sharedVar[(itemX + 1) % 16]);
    }
}
```

OKL

## Exclusive memory

```
for(int groupX = 0; groupX < xGroups; ++groupX; outer0){ // Work-group implicit loops
    exclusive int exclusiveVar, exclusiveArray[10];

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        exclusiveVar = itemX; // Pre-fetch
    }

    // Auto-insert [barrier(localMemFence);]

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        int i = exclusiveVar; // Use pre-fetched data
    }
}
```

OKL

# OKL: OCCA Kernel Language

## Shared memory

```
for(int groupX = 0; groupX < xGroups; ++groupX; outer0){ // Work-group implicit loops
    shared int sharedVar[16];

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        sharedVar[itemX] = itemX;
    }

    // Auto-insert [barrier(localMemFence);]

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        int i = (sharedVar[itemX] + sharedVar[(itemX + 1) % 16]);
    }
}
```

OKL

## Exclusive memory

```
exclusiveVar = 0;
exclusiveVar = 1;
exclusiveVar = 2;
.
.
.
for(int groupX = 0; groupX < xGroups; ++groupX; outer0){ // Work-group implicit loops
    exclusive int exclusiveVar, exclusiveArray[10];

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        exclusiveVar = itemX; // Pre-fetch
    }

    // Auto-insert [barrier(localMemFence);]

    for(int itemX = 0; itemX < 16; ++ itemX; inner0){ // Work-item implicit loops
        int i = exclusiveVar; // Use pre-fetched data
    }
}
```

OKL

# OFL: OCCA Fortran Language

## Description

- Translates to OKL and then to OCCA IR with code transformations
- Parallel loops are explicit through the **inner** and **outer** DO-labels

```
kernel subroutine kernelName(...)
...

DO groupY = 1, yGroups, outer1
  DO groupX = 1, xGroups, outer0 // Work-group implicit loops
    DO itemY = 1, yItems, inner1
      DO itemX = 1, xItems, inner0 // Work-item implicit loops
        // GPU Kernel Scope
      END DO
    END DO
  END DO
END DO

...
end subroutine kernelName
```

OFL

## Shared and exclusive memory

```
integer(4), shared    :: sharedVar(16,30)
integer(4), exclusive :: exclusiveVar, exclusiveArray(10)
```

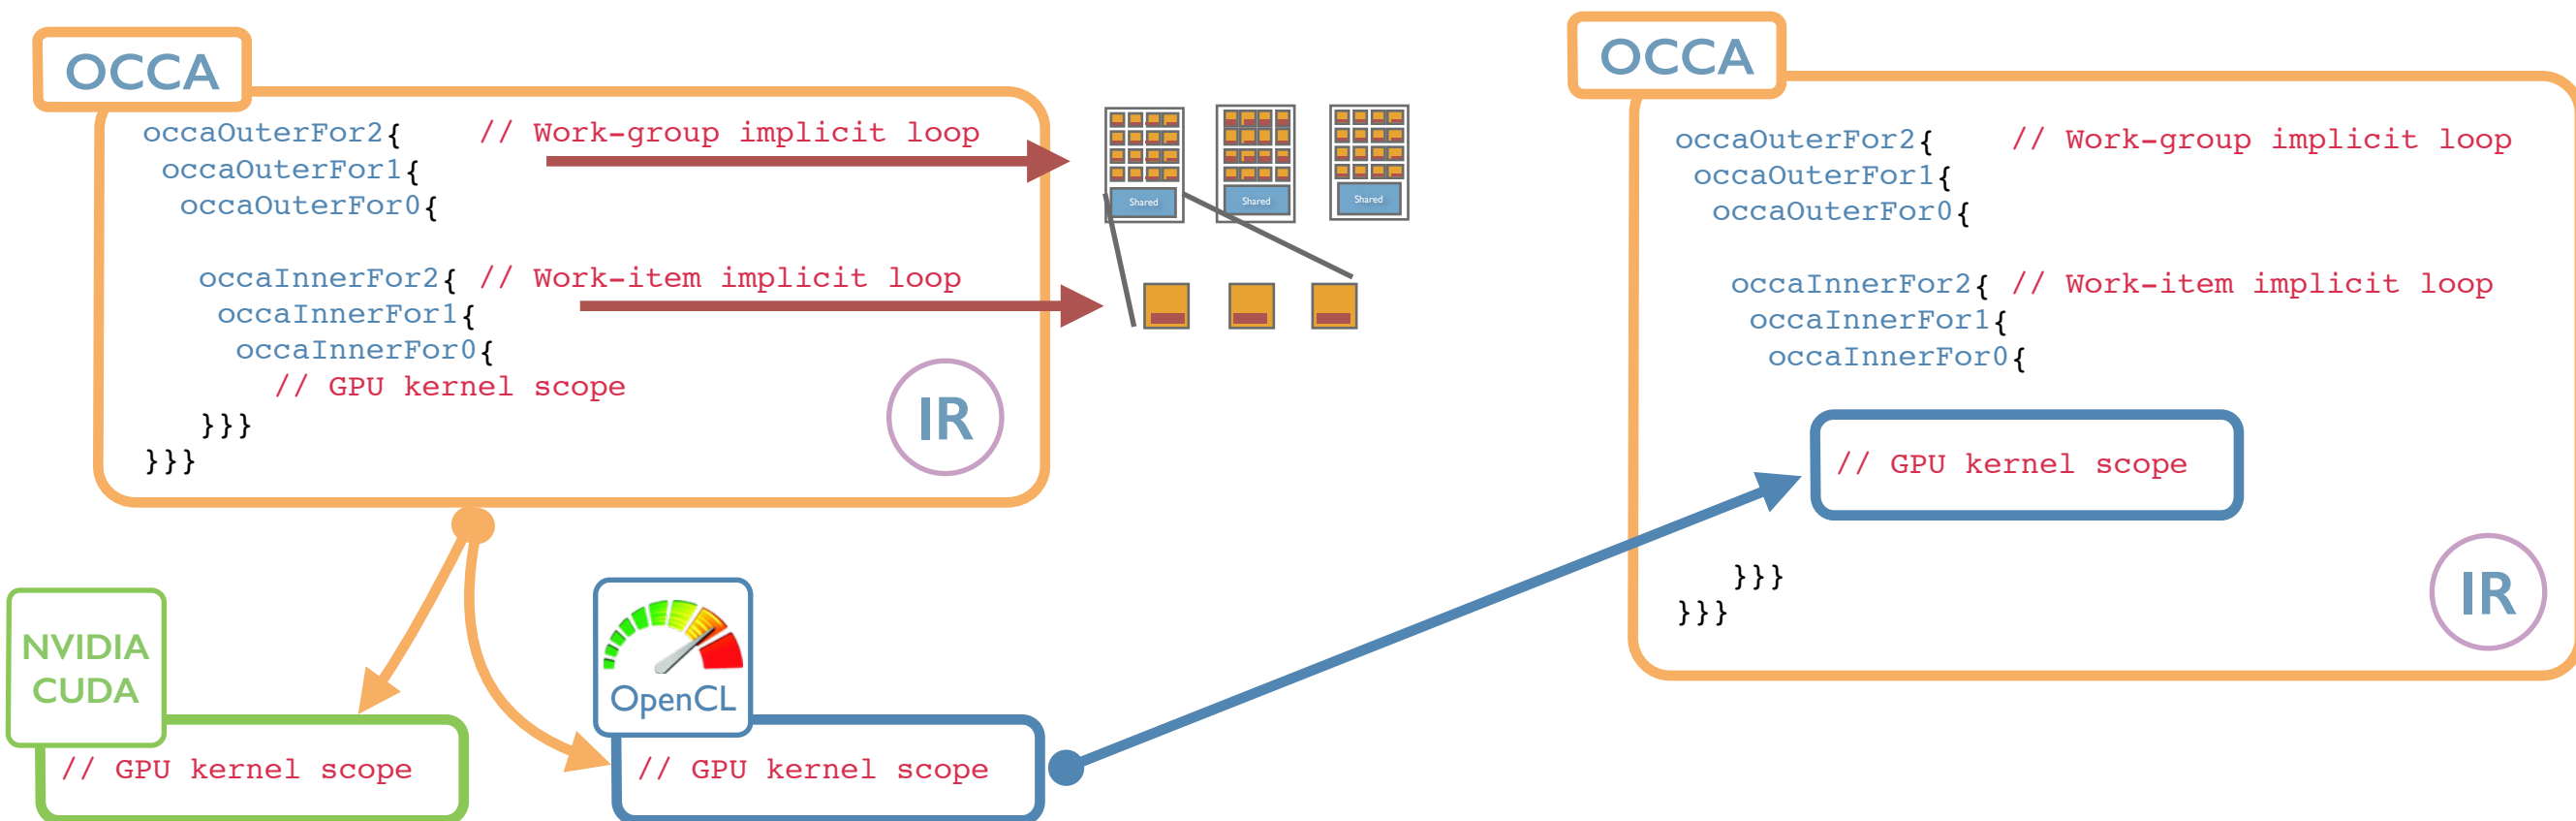
OFL



# OpenCL/CUDA to OCCA IR

## Description


- Parser can translate OpenCL/CUDA kernels to OCCA IR\*
- Although OCCA IR was derived from the GPU model, there are complexities




Since we derived OCCA IR from the GPU model, the inverse should be easy ... right?

# OCCA2: apps & benchmarks


OCCA apps can perform close to or exceed native apps across platforms.



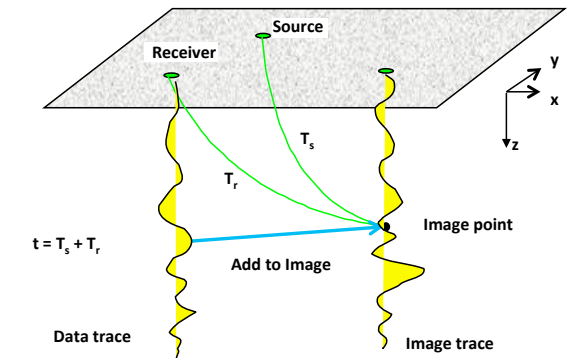
FDTD RTM: wave equation (+15-33%)




Lattice Boltzmann for Core Sample Analysis



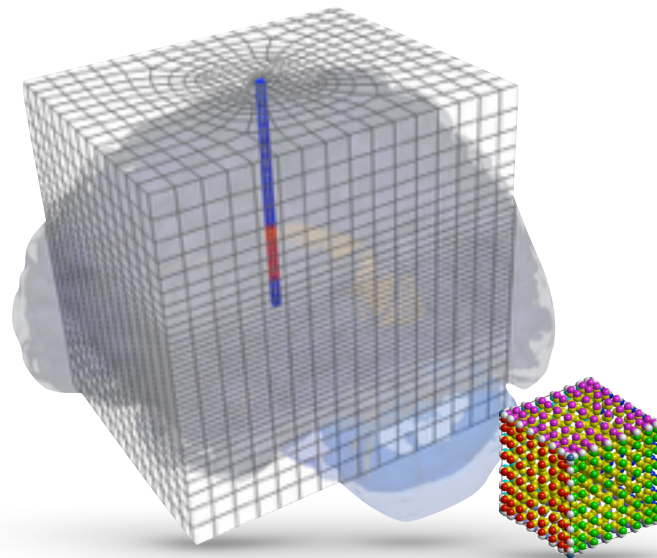
DG for seismic wave simulations



Kirchhoff migration (kernel up to 2x faster).



ALMOND: algebraic multigrid library



MDACC: FEM model of laser tumor ablation



DG compressible Navier-Stokes solver

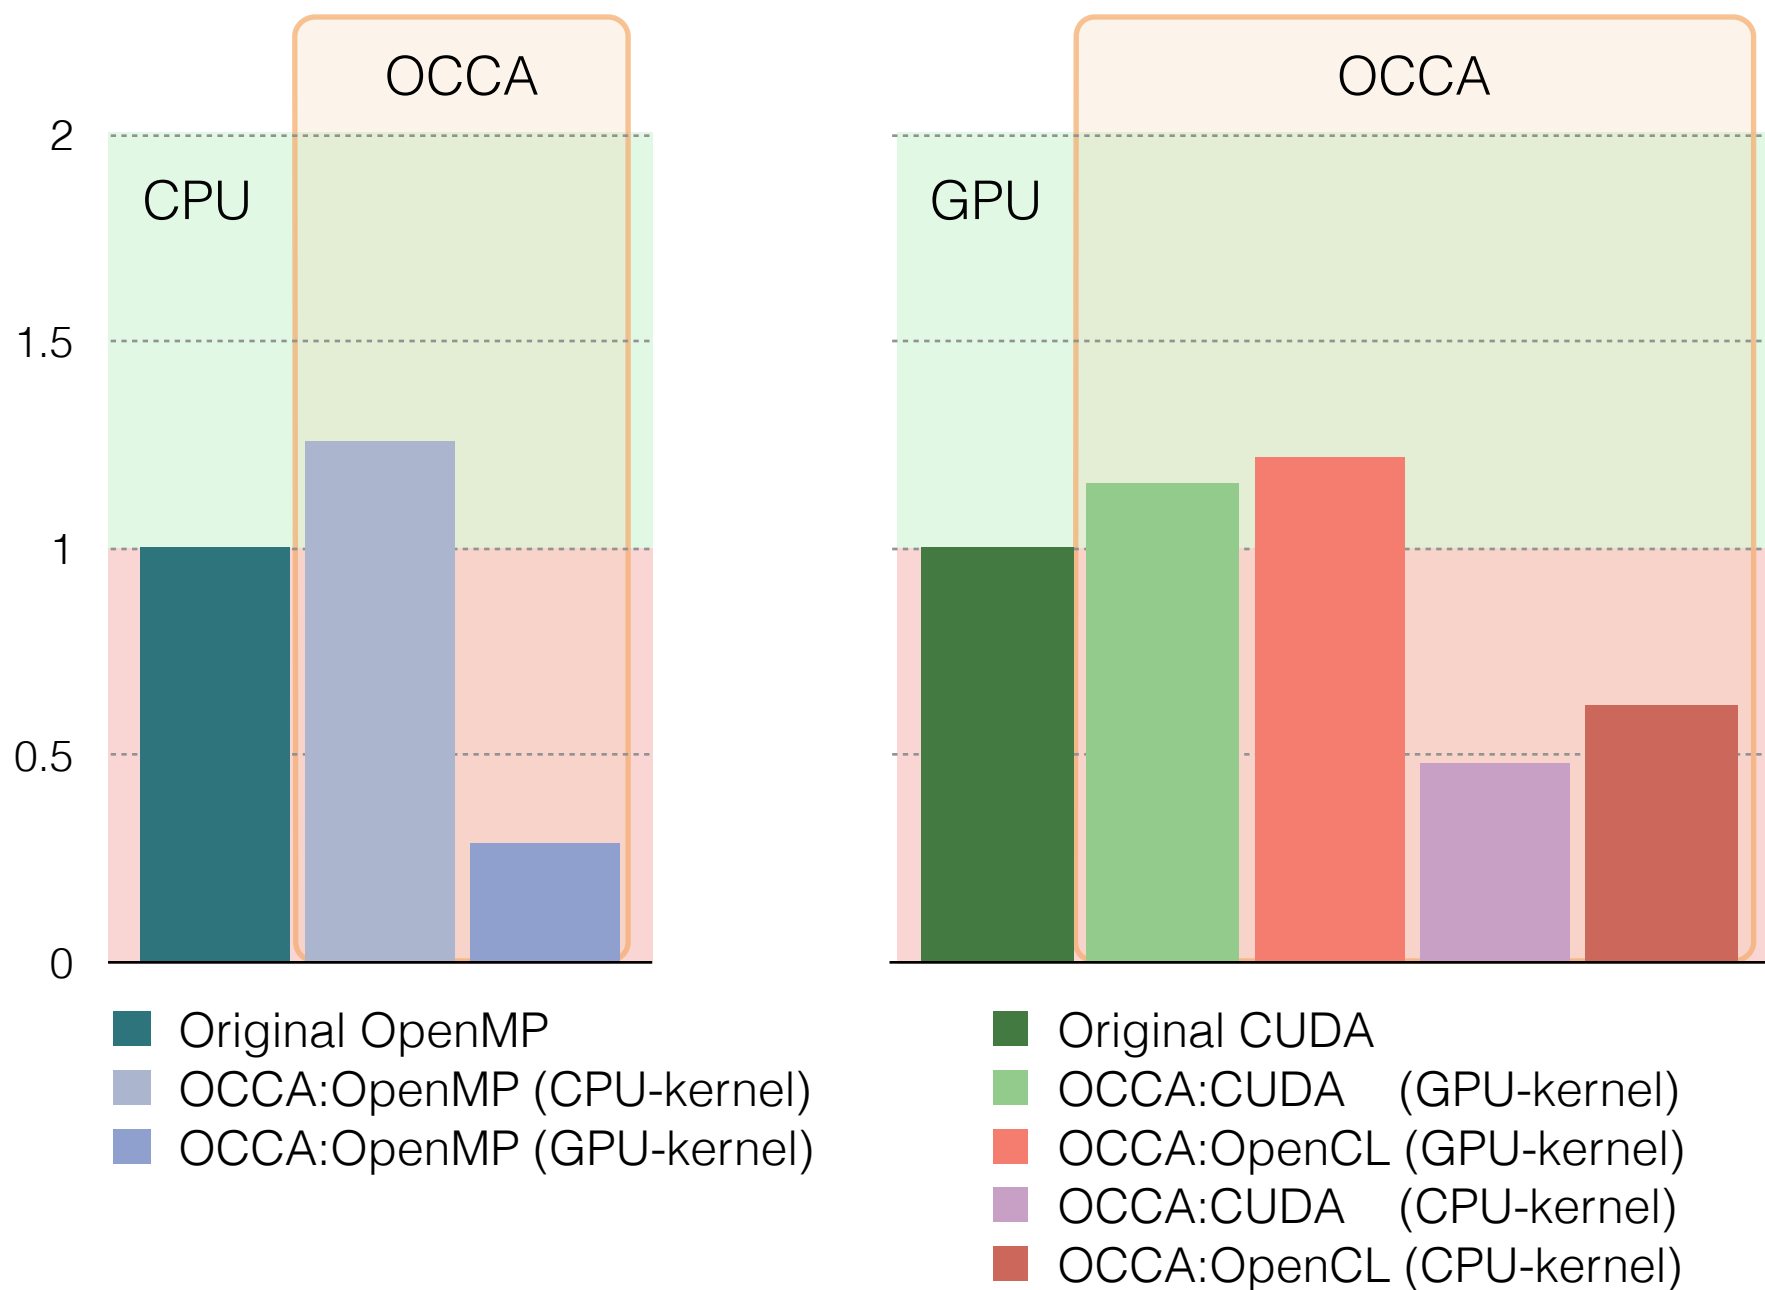


DG shallow-water & 3D ocean modeling

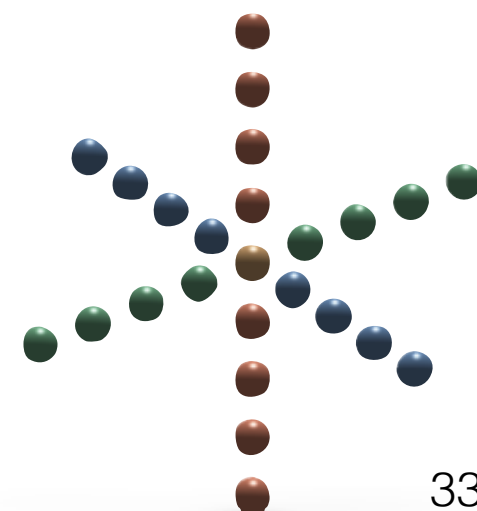
+ more applications

# OCCA2: apps & benchmarks

High-order finite difference for RTM

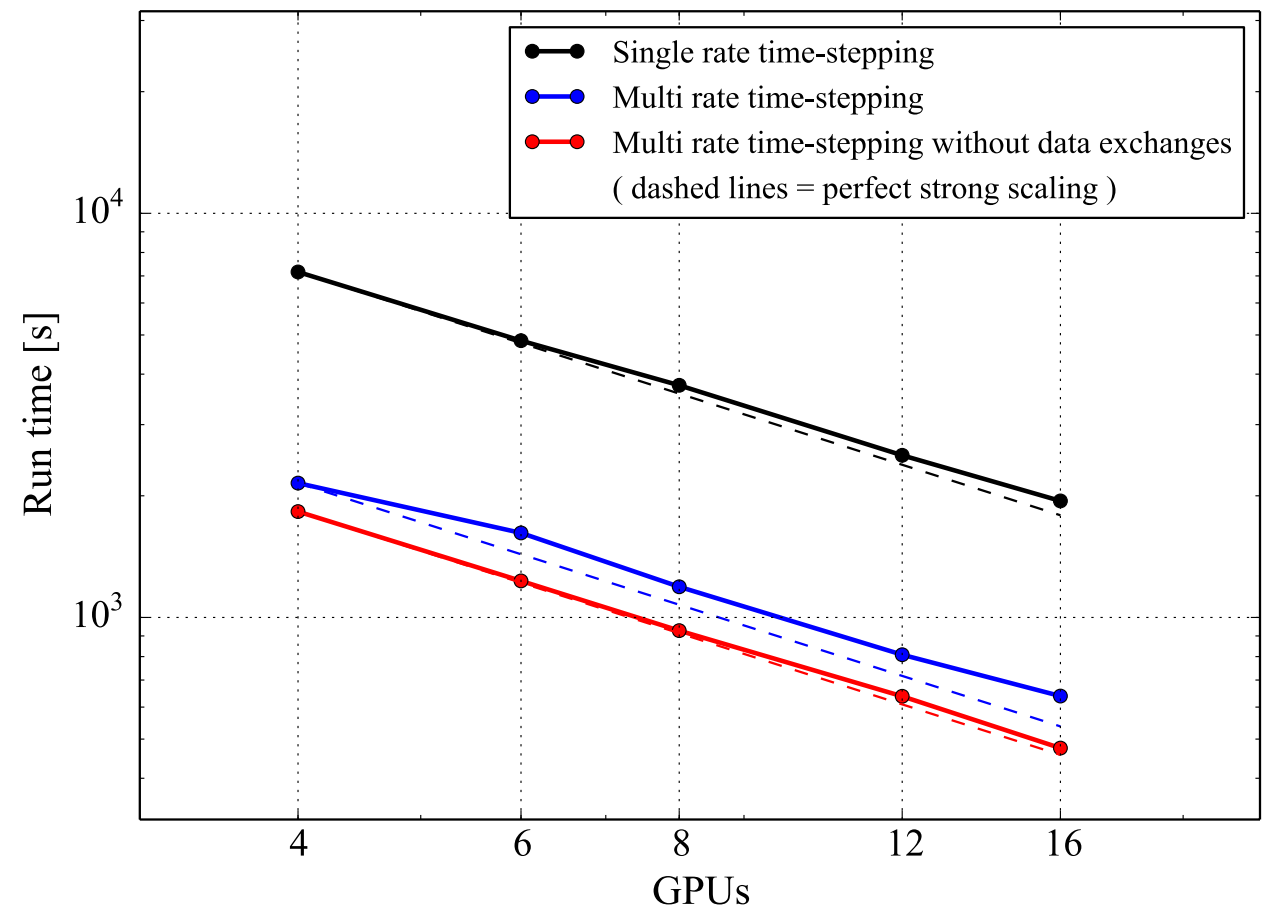



OpenMP : Intel Xeon CPU E5-2640  
OpenCL/CUDA : NVIDIA Tesla K10

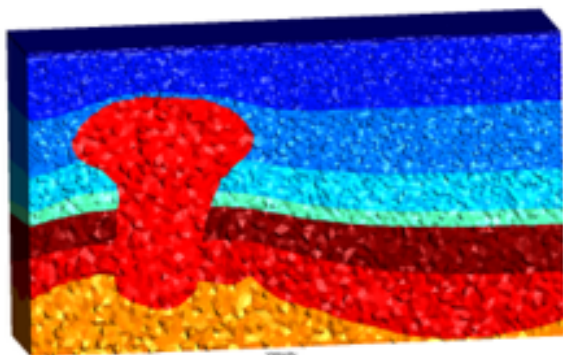


# OCCA2: apps & benchmarks

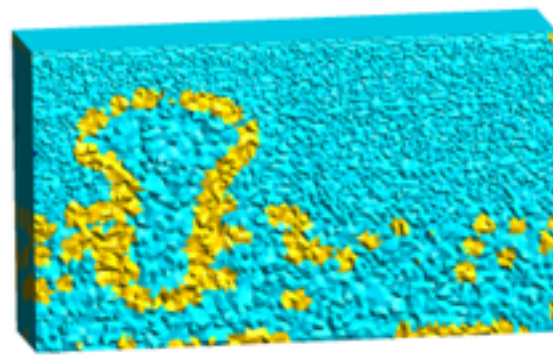
High-order discontinuous Galerkin for RTM




Discretization



Multi-rate scheme



Distributed



# OCCA2: apps & benchmarks

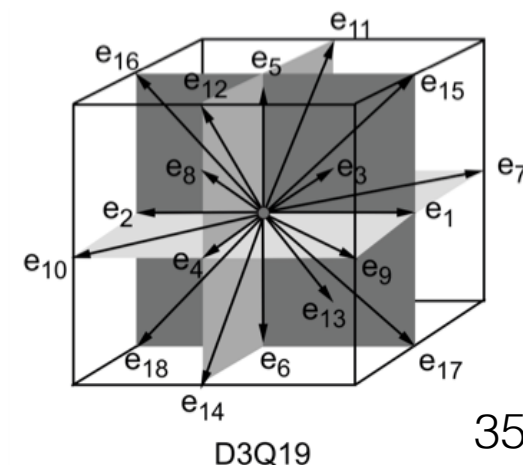
Lattice Boltzmann Method in Core Sample Analysis

Comparison across platforms (Normalized with original code)

	API Mode	Device	Model	Wall Clock	BW (GB/s)	Speedup
	Ref dense code [-O3 in gcc 4.8]	CPU 1-core	Intel i7-5960X	1290	—	x 1
OCCA	OpenMP	CPU	Intel i7-5960X	11.12	22	x 116
	OpenCL: Intel	CPU	Intel i7-5960X	11.18	22	x 115
	OpenCL: AMD	GPU	AMD 7990	1.39	176	x 928
	OpenCL: NVIDIA	GPU	GTX 980	1.25	196	x 1032
	CUDA: NVIDIA	GPU	GTX 980	1.20	205	x 1075

Comparison across platforms (Normalized with OCCA::OpenMP)

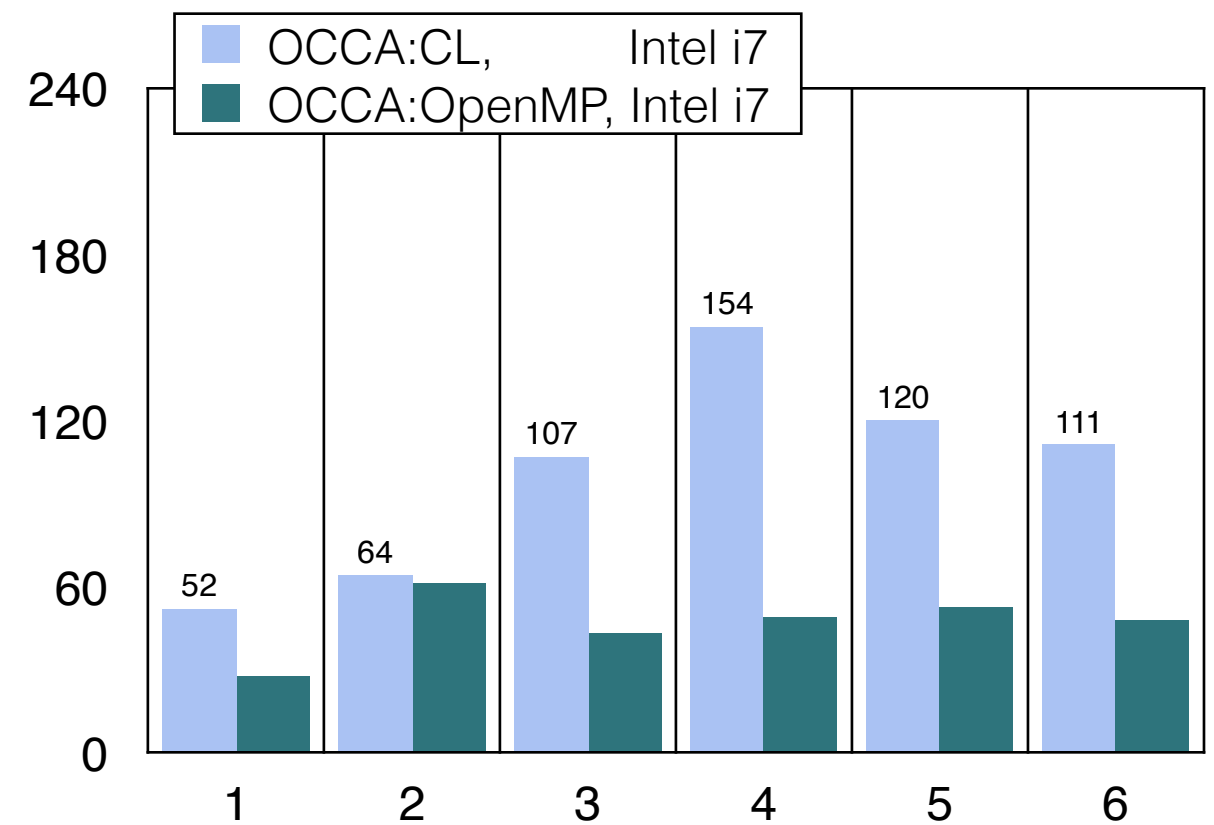

	API Mode	Device	Model	Wall Clock	BW (GB/s)	Speedup
OCCA	OpenMP	CPU	Intel i7-5960X	11.12	22	x 1.0
	OpenCL: Intel	CPU	Intel i7-5960X	11.18	22	x 1.0
	OpenCL: AMD	GPU	AMD 7990	1.39	176	x 8.0
	OpenCL: NVIDIA	GPU	GTX 980	1.25	196	x 8.9
	CUDA: NVIDIA	GPU	GTX 980	1.20	205	x 9.3





# OCCA2: apps & benchmarks

Discontinuous Galerkin for shallow water equations



Polynomial Order	Compute-Time vs Real-Time
1	x650
2	x208
3	x95
4	x47

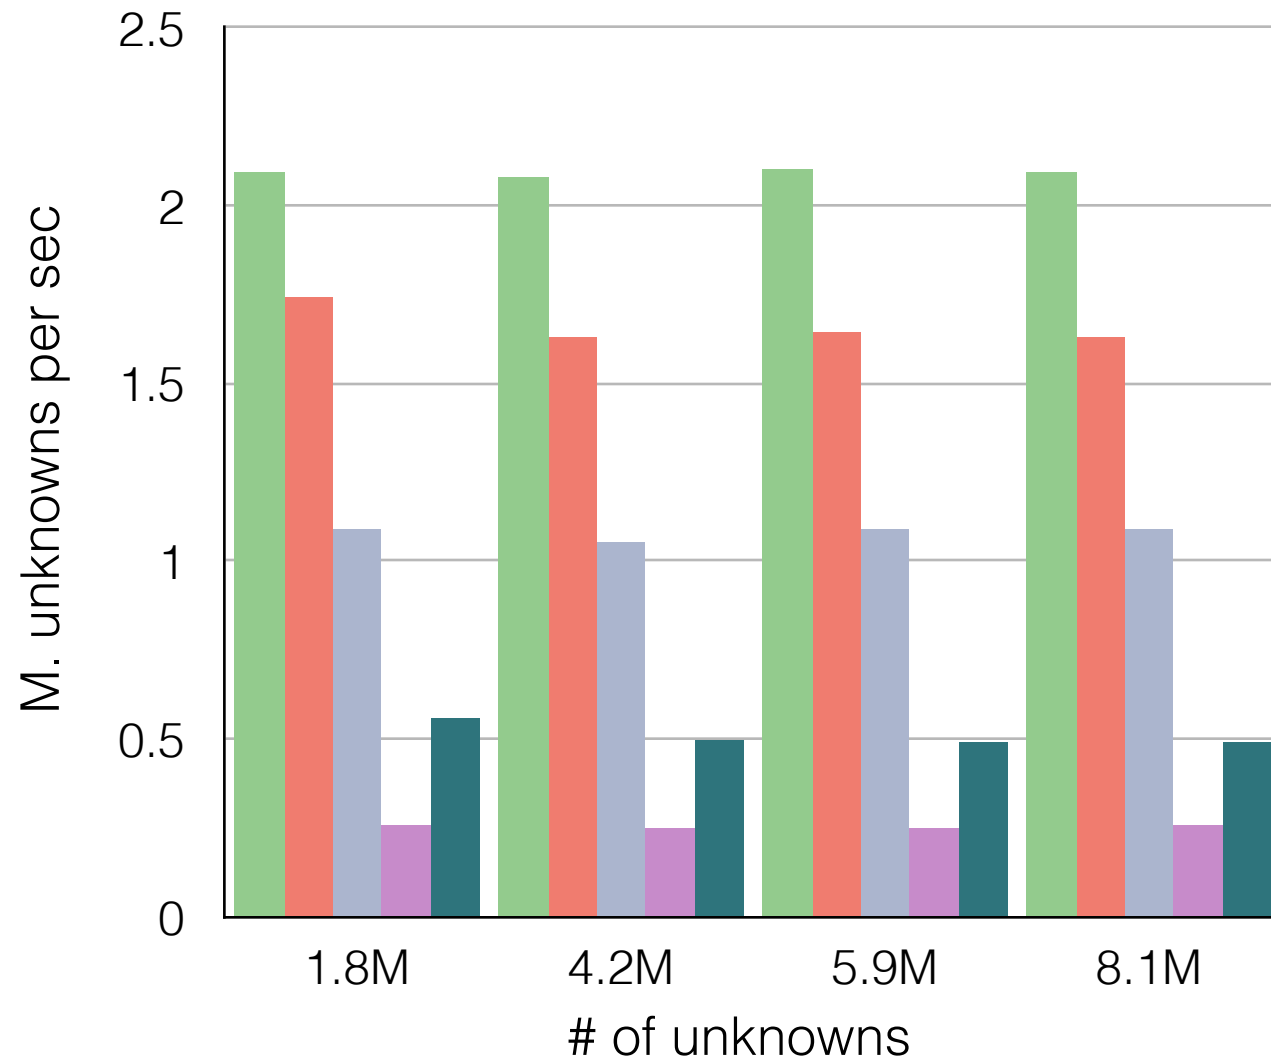




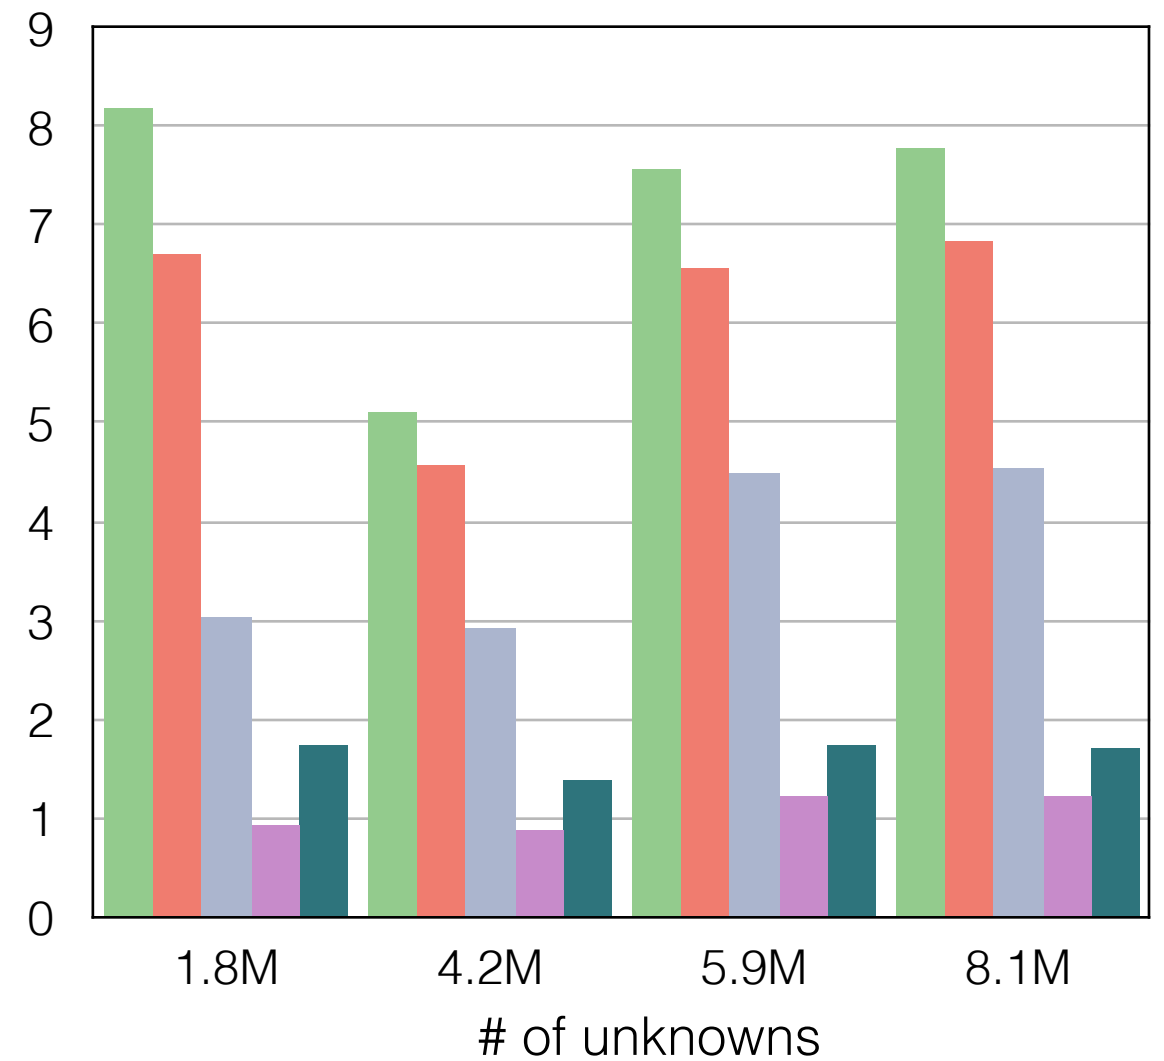
# OCCA2: apps & benchmarks

Algebraic multigrid for elliptic problems

Setup Time




Solve Time



CUDA on Titan  
OpenCL on Intel i7


OpenCL on Titan  
OpenMP on Intel i7

OpenCL on Tahiti



# OCCA2: apps & benchmarks

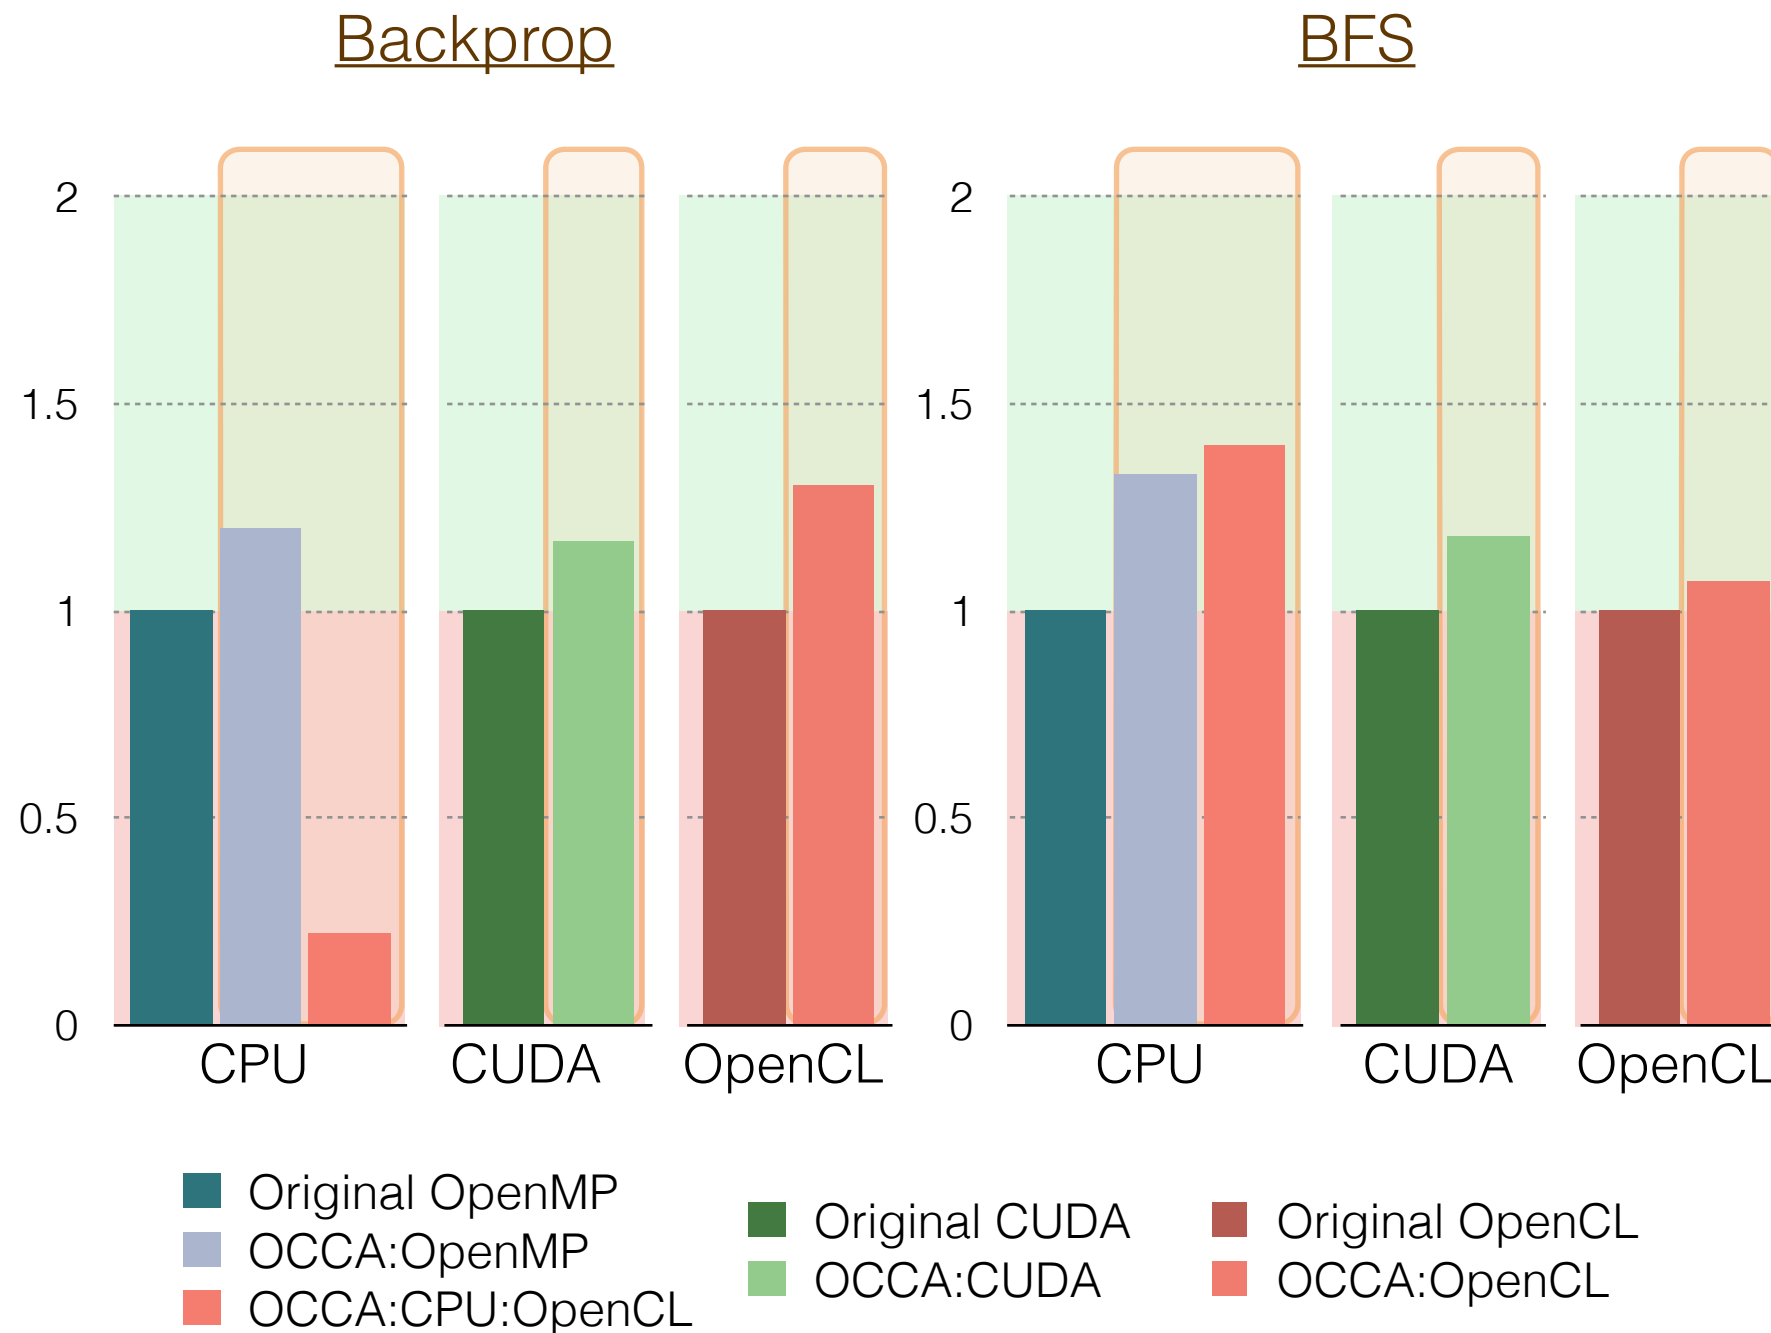
Monte Carlo for neutronics  
Collaborations with Argonne National Lab



OpenMP : Intel Xeon CPU E5-2650  
OpenCL/CUDA : NVIDIA Tesla K20c

# OCCA2: apps & benchmarks

Two of our ported Rodinia benchmarks, based on the “11 Dwarves”



OpenMP : Intel Xeon CPU E5-2650  
OpenCL/CUDA : NVIDIA Tesla K20c

# Live Demo

## Installation

- Download OCCA and template code for this session at:

```
git clone https://github.com/tcew/OCCA2.git  
git clone https://github.com/tcew/OG15.git
```

- Setup environment variables and compile

```
cd OCCA2  
export OCCA_DIR=$PWD  
export CXX=clang++ # or any compiler you want  
export PATH=$PATH:$OCCA_DIR/bin  
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$OCCA_DIR/lib  
make -j
```

- Try it out

```
occainfo # Displays available devices  
cd examples/addVectors  
make -j  
./main
```


# Interactive Demo

## Description

- Go to the previously downloaded template code:

```
cd $OCCA_DIR/../OG15 # If you downloaded it next to OCCA  
cd matrixTranspose
```

- This demo will let you (the attendees) try coding an OKL kernel (or OFL)
- Translate the serial matrix transpose seen in the code into an OKL/OFL kernel



A few of us ^ will be going around, feel free to raise your hand for questions ☺